

UNIVERSITY OF LONDON  
IMPERIAL COLLEGE OF SCIENCE, TECHNOLOGY AND MEDICINE

EXAMINATIONS 2000

BEng Honours Degree in Computing Part II  
MEng Honours Degrees in Computing Part II  
for Internal Students of the Imperial College of Science, Technology and Medicine

*This paper is also taken for the relevant examinations for the  
Associateship of the City and Guilds of London Institute*

PAPER C211

OPERATING SYSTEMS II

Thursday 11 May 2000, 14:00  
Duration: 90 minutes  
(Reading time 5 minutes)

*Answer THREE questions*

Paper contains 4 questions

- 1
  - a Is the code in the OS that does scheduling part of a task or not? Explain your answer.
  - b Explain how the scheduler works, how its gets activated, and what activates it. Give a diagram that describes the Round Robin Scheduling Algorithm.
  - c Give the *process state diagram*. Give an extension that deals with the notion of a process being *swapped out*, i.e., the process descriptor exists, but no longer has its code in memory.
  - d Describe in a few lines the effect of the Unix system calls `fork` and `exec` and how they interact.

Try to be concise for each case. 100 words per part should be the maximum.

The four parts carry, respectively, 15%, 30%, 30%, and 25% of the marks.

- 2
- a Describe, using a diagram, in which way the blocks used for a file are administered within a MINIX i-node. Give at least three other items of data contained in an i-node located on disk. Give also the extra items of data that are kept in i-nodes for open files.
  - b Give the names of at least four disk scheduling algorithms and describe them in a few words.
  - c Describe via a picture the tables used in the typical file sharing model. Give a short description of what is kept for each entry in the tables involved.
  - d Discuss the UNIX/MINIX file sharing model. In your answer, describe the data management by the OS when a file is opened; make sure to discuss the various file tables involved (you can ignore the i-node handling).

The four parts carry, respectively, 25%, 20%, 30%, and 25% of the marks.

- 3 a Describe each of the following access control methods, and explain how they are used (for each use no more than 30 words)
- Protection Domain
  - Access Matrix
  - Access Control List
  - Global Table
  - Capability
- b Discuss how capabilities can be revoked in a distributed system.
- c The following access matrix is used to control access to resources in a system

	File1	File2	File3	Printer	D1	D2	D3	D4
D1	<i>read, write+</i>					<i>enter</i>		
D2	<i>read</i>	<i>write</i>	<i>owner</i>					
D3	<i>read</i>	<i>read</i>	<i>exec</i>	<i>owner, print</i>				<i>enter</i>
D4		<i>exec</i>	<i>read</i>	<i>print*</i>		<i>control</i>		

The following are members of the domains:

- D1 - Charles  
D2 - Mary, Fred  
D3 - Fred, James, Joan  
D4 - Pat

Give your answers to the following questions, and explain how you are able to derive these answers:

- i) Who can write to File1?
- ii) How can Charles print to the Printer?
- iii) Identify all resources and operations accessible to Fred.
- iv) Can Mary lose the right to write to File2?

The three parts carry, respectively, 30%, 40%, and 30% of the marks.

- 4    a    Give, in a picture, the structure of process queues in MINIX, and discuss how the scheduler deals with them.
- b    In MINIX, both asynchronous send and synchronous send are used for message passing. A message sent by a process gets sent asynchronously if the receiver is not waiting for a message, and has no message pending; one message can be stored for each process. Otherwise, the sending process blocks when sending to a process that has already a message pending. Explain how messages sent within MINIX still can get lost.
- c    In the process administration of the Unix-like operating system MINIX, process 0 (zero) corresponds to the pseudo-process HARDWARE, that serves as the source of the messages which are sent as the direct result of an interrupt. Although this pseudo-process has a number, it does not have a slot in the Process Table (hence the predicate 'pseudo'). Explain this apparant contradiction.
- d    In MINIX, system calls are implemented using `SendRec`, forcing the user process to explicitly wait for the result from the requested action to arrive. Kernel processes (or tasks), however, can send messages without having to wait explicitly for an answer. Explain how the implementation of communication between user processes and the kernel avoids getting into a deadlock, and how the specific organisation of the message passing between tasks in MINIX avoids deadlocks.

The four parts carry, respectively, 20%, 20%, 20% and 40% of the marks.

*End of Paper*