

Paper Number(s): **EE1-9A**

IMPERIAL COLLEGE LONDON

DEPARTMENT OF ELECTRICAL AND ELECTRONIC ENGINEERING
EXAMINATIONS 2011

ISE Part I: MEng, BEng and ACGI

Corrected Copy Q2 (b)

INTRODUCTION TO COMPUTER ARCHITECTURE AND SYSTEMS (PART A)

Monday, 6 June 2:00 pm

Time allowed: 1:30 hours

There are THREE questions on this paper.

Answer ALL questions.

Question 1 carries 40% of the marks. Questions 2 and 3 carry equal marks (30% each).

Any special instructions for invigilators and information for candidates are on page 1.

Examiners responsible:

First Marker(s): Clarke, T.
Second Marker(s): Demiris, Y.

Special instructions for invigilators

The sheet Exam Notes 2011 should be distributed with the Examination Paper.

Special instructions for students

The prefix &, or suffix $_{(16)}$, introduces a hexadecimal number, e.g: &1C0, 1C0 $_{(16)}$.

Unless otherwise specified negative numbers are represented in two's complement.

Unless otherwise specified machine addressing is little-endian.

The sheet Exam Notes 2011, as published on the course web pages, is provided and contains reference material.

Answer ALL the questions.

The Questions

1.

a)

- (i) Calculate the 32 bit IEEE-754 representation of $-31.25_{(10)}$, writing your answer as 8 hexadecimal digits.
- (ii) The value &FF is stored in an N bit register. State the two's complement signed value of the register in the two cases $N=16$ and $N=8$.

[8]

- b) A write-through direct mapped cache has line (block) length of 16 32-bit words, and contains of 512 lines. State the total data storage size of the cache (not counting tag memory) in bytes, and the tag, index and select fields for this cache in a 32 bit ARM address.

[8]

- c) Write a fragment of ARM assembly code which implements in as few instructions as possible:

$R3 := R1 + R2 - R3 + 111$

[8]

- d) For each of the ARM instructions below, summarise the datapath transfer that occurs by giving one of the 4 options in Figure 1.1. In cases A & B only, specify what is the destination register. Note that in the ARM architecture PC is a data register.

- (i) **MOV R0,R1**
- (ii) **STRB R0, [R1]**
- (iii) **LDR R0, [R1,R2]**
- (iv) **B START**

[8]

- e) State single ARM instructions which perform the following integer operations on unsigned 32 bit operands, giving a 32 bit result, rounding down, and ignoring overflow.

- (i) $R10 := R10 + R9 * 4096$
- (ii) $R3 := R3 + R3 / 8$

[8]

A	data register(s) => data register
B	memory unit location => data register
C	data register => memory unit location
D	no datapath transfer

Figure 1.1. Datapath transfers

2. Each code fragment (a) - (c) below executes with all condition codes and registers initially 0, and memory locations as in *Figure 2.1*. State the values of R0-R3, and the condition codes, after execution of the code fragment. Write your answers using as a template a copy of the table in *Figure 2.3*, deleting the example row labelled (x) which indicates the required format of your answer. Each answer must be written in hexadecimal except the condition codes, which must be in binary, as indicated in row (x).

a) Code as in *Figure 2.2a*.

[10]

b) Code as in *Figure 2.2b*.

[10]

c) Code as in *Figure 2.2c*. Note that code fragment execution terminates directly after the BGE instruction is condition false executed.

[10]

Location (word)	Value
&100	&04030201
&104	&08070605
> &104	&0

Figure 2.1. Memory locations

```
MOV R0, #1
MVN R1, R0
ADDS R2, R0, R1
SBC R3, R0, R1
ADC R0, R0, R0
```

(a)

```
MOV R10, #100
LDR R0, [R10], #4
LDR R1, [R10], #4
ORR R2, R0, R1
ANDS R3, R0, R1
```

(b)

```
MOV R2, #3
MOV R1, #&100
L1 LDRB R0, [R1, #1]!
ORR R3, R0, R3, lsl #8
SUBS R2, R2, #1
BGE L1
```

(c)

Figure 2.2. Code fragments

	R0	R1	R2	R3	NZCV
(x)	&0	&1020	&FFFFFFFF	&C	0110
(a)					
(b)					
(c)					

Figure 2.3. Template for answers

3.

- a) Explain the precise operation of the **LDMED** & **STMED** instructions in the subroutine **PARITY** displayed in *Figure 3.1*. Indicate the function of **R13**, and how the transfers in these instructions implement the subroutine.

[6]

- b) The code in *Figure 3.2* which calls the **PARITY** subroutine is executed from **START** to **FINISH** (not including the unspecified instruction after label **FINISH**). Using the instruction timings in the exam notes, work out the total number of cycles used during this execution where:

- (i) Conditional branch **BEQ** is condition true executed.
- (ii) Conditional branch **BEQ** is condition false executed.

Explain, with reference the ARM hardware, the given timing for the **STMED** instruction.

[8]

- c) By computing the value of **R2** LSB throughout execution of the **PARITY** subroutine, or otherwise, show how the **Z** status bit on subroutine exit relates to the bits of **R1** on entry.

[8]

- d) The computation from **START** to **FINISH** can be optimised as follows:
- Use **R1** to calculate the **PARITY** result, overwriting **R1** in the process. Other than the change to **R1** *the subroutine must have the same specification as before*.
 - Eliminate the conditional branch instruction in the calling code.

Rewrite the code in *Figure 3.1* and *Figure 3.2* with these changes in such a way that the value of **R1** at **FINISH** remains the same as in the original code.

[8]

```

; R1: input
; Z status bit: output
; no register is changed.
PARITY STMED    R13!, {R2,R14}
      MOV      R2, R1
A      EOR      R2, R2, R2, ror #1
      EOR      R2, R2, R2, ror #2
      EOR      R2, R2, R2, ror #4
B      ANDS     R2, R2, #1
      LDMED     R13!, {R2,R15}

```

Figure 3.1. PARITY subroutine


```

START  BL  PARITY
      BEQ  FINISH
      ORR  R1, R1, #&100
FINISH

```

Figure 3.2. Code calling the PARITY subroutine

Data processing (ADD,SUB,AND,CMP,MOV, etc)



The second operand, Op2, is either a constant C or register Rm

Assume Shift=0, Rot=0, for unshifted Rm or immediate C

C	1 => carry
V	1 => signed overflow
N	1 => negative
Z	1 => zero

Op-codes	
AND	EQ,NE,... => Condition
ANDEQ	S => set status on result
ANDS	note position of S
ANDEQS	

Op	Assembly	Operation	Pseudocode
0000	AND Rd,Rn,op2	Bitwise logical AND	Rd := Rn AND op2
0001	EOR Rd,Rn,op2	Bitwise logical XOR	Rd := Rn XOR op2
0010	SUB Rd, Rn, op2	Subtract	Rd := Rn - op2
0011	RSB Rd, Rn, op2	Reverse subtract	Rd := op2 - Rn
0100	ADD Rd,Rn,op2	Add	Rd := Rn + op2
0101	ADC Rd,Rn,op2	Add with carry	Rd := Rn + op2 + C
0110	SBC Rd, Rn, op2	Subtract with carry	Rd := Rn - op2 - C
0111	RSC Rd, Rn, op2	Reverse sub with C	Rd := op2 - Rn + C
1000	TST Rn, op2	set NZ on AND	Rn AND op2
1001	TEQ Rn, op2	set NZ on EOR	Rn EOR op2
1010	CMP Rn, op2	set NZCV on -	Rn - op2
1011	CMN Rn, op2	set NZCV on +	Rn + op2
1100	ORR Rd,Rn,op2	Bitwise logical OR	Rd := Rn OR op2
1101	MOV Rd, op2	Move	Rd := op2
1110	BIC Rd,Rn,op2	Bitwise clear	Rd := Rn AND NOT op2
1111	MVN Rd,op2	Bitwise move invert	Rd := NOT op2

Data Processing Op2

ADD r0, r1, op2
MOV r0, op2

Examples

ADD r0, r1, r2
MOV r0, #1
CMP r0, #1
EOR r0, r1, r2, lsr #10
RSB r0, r1, r2, asr r3

Op2	Conditions	Notes
Rm		r15=pc, r14=lr, r13=sp
#imm	imm = s rotate 2r (0 ≤ s ≤ 255, 0 ≤ r ≤ 15)	Assembler will translate negative values changing op-code as necessary Assembler will work out rotate if it exists
Rm, shift #s Rm, rrx #1	(1 ≤ s ≤ 31) shift => lsr,lsr,asr,asr,ror	rrx always writes carry ror writes carry if S=1 shifts do not write carry
Rm, shift Rs	shift => lsr,lsr,asr,asr,ror	shift by register value (takes 2 cycles)

Multiply in detail

- MUL,MLA were the original (32 bit LSW result) instructions
 - Why does it not matter whether they are signed or unsigned?
- Later architectures added 64 bit results


Register operands only
No constants, no shifts

ARM3 and above

		NB d & m must be different for MUL, MLA	
MUL	rd, rm, rs	multiply (32 bit)	Rd := (Rm*Rs)[31:0]
MLA	rd,rm,rs,rn	multiply-acc (32 bit)	Rd := (Rm*Rs)[31:0] + Rn
UMULL	rh, rl, rm, rs	unsigned multiply	(Rh:Rl) := Rm*Rs
UMLAL	rh, rl, rm, rs	unsigned multiply-acc	(Rh:Rl) := (Rh:Rl) + Rm*Rs
SMULL	rh,rl,rm,rs	signed multiply	(Rh:Rl) := Rm*Rs
SMLAL	rh,rl,rm,rs	signed multiply-acc	(Rh:Rl) := (Rh:Rl) + Rm*Rs

ARM7DM core and above (64 bit multiply result)

Data transfer (to or from memory LDR,STR)



Bit in word	0	1
P	use base register addressing [Rn]	use indexed or offset address [Rn+Rm], [Rn+S]
U	subtract offset [Rn-S]	add offset [Rn+S]
B	Word	Byte
W	leave Rn unchanged if P=1	write indexed or offset address back into Rn if P=1
L	Store	Load

NB - P=0, W=1 is not allowed

If P=0, W=0, write offset address back into Rn. For no change to Rn & no offset use P=0,W=0,S=0

Data Transfer Instructions

LDR	load word
STR	store word
LDRB	load byte
STRB	store byte
LDREQB	NB B is after EQ condition
STREQB	

LDR	r0, [r1]	register-indirect addressing
LDR	r0, [r1, #offset]	pre-indexed addressing (base + offset)
LDR	r0, [r1, #offset]!	pre-indexed, auto-indexing (base + offset + writeback)
LDR	r0, [r1], #offset	post-indexed, auto-indexing (change Rn after)
LDR	r0, [r1, r2]	register-indexed addressing (base + reg)
LDR	r0, [r1, r2, lsr #shift]	scaled register-indexed addressing (base + reg * 2 ^{shift})
LDR	r0, address_label	PC relative addressing (pc+8 is read, offset calculated)
ADR	r0, address_label	load PC relative address (pc+8 is read, offset calculated)

LDMED r13!, {r0-r4,r6,r6}; !=> write-back to address register
 STMFA r13, {r2}; no write-back
 STMEQIB r2!, {r5-r12}; note position of EQ or other condition
higher reg nos go to/from higher mem addresses
 [E|F][A|D] Empty|Full, Ascending|Descending
 [I|D][A|B] Increment|Decrement, After|Before

Name	Stack	Other
pre-increment load	LDMED	LDMIB
post-increment load	LDMFD	LDMIA
pre-decrement load	LDMEA	LDMDB
post-decrement load	LDMFA	LDMDA
pre-increment store	STMFA	STMIB
post-increment store	STMEA	STMIA
pre-decrement store	STMFD	STMDB
post-decrement store	STMED	STMDA

Instruction Timing

Exact instruction timing is very complex and depends in general on memory cycle times which are system dependent. The table below gives an approximate guide.

Instruction	Typical execution time (cycles)
Any instruction, with condition false	1
data processing (except register-valued shifts)	1 (+3 if Rd = R15)
data processing (register-valued shifts): MOV R1, R2, lsl R3	2 (+3 if Rd = R15)
LDR, LDRB, STR, STRB	4 (+3 more if Rd = R15)
LDM (n registers)	n+3 (+3 more if Rd = R15)
STM (n registers)	n+3
B, BL	4
Multiply	7-14

$$x = (-1)^s 2^{(e-127)} 1.f$$

IEEE 754



Shadow registers

- FIQ mode: R8 - R14 shadowed
- IRQ, SVC, abort, UND modes: R13 - R14 shadowed
- Return from interrupt: set status bits to restore CPSR, so use SUBS to set PC equal to stored return address with offset & restore CPSR.
- R13 is SP

BL, SWI	R14 = Return address
IRQ, FIQ, UND	R14 = Return address + 4
Abort	R14 = Return address + 8

Exception	Mode	Vector address
Reset	SVC	0x00000000
Undefined instruction	UND	0x00000004
Software interrupt (SWI)	SVC	0x00000008
Prefetch abort (instruction fetch memory fault)	Abort	0x0000000C
Data abort (data access memory fault)	Abort	0x00000010
IRQ (normal interrupt)	IRQ	0x00000018
FIQ (fast interrupt)	FIQ	0x0000001C

Cond		Condition	Status Bits
0000	EQ	Equal	Z set
0001	NE	Not equal	Z clear
0010	CS/HS	Unsigned ≥ (High or Same)	C set
0011	CC/LO	Unsigned < (Low)	C clear
0100	MI	Minus (negative)	N set
0101	PL	Plus (positive or 0)	N clear
0110	VS	Signed overflow	V set
0111	VC	No signed overflow	V clear
1000	HI	Unsigned > (High)	C set and Z clear
1001	LS	Unsigned ≤ (Low or Same)	C clear OR Z set
1010	GE	Signed ≥	N equals V
1011	LT	Signed <	N is not equal to V
1100	GT	Signed >	Z clear and N equals V
1101	LE	Signed ≤	Z set and N not equal to V
1110	AL	Always	any
1111	NV	Never (do not use)	none

Machine Instruction Overview (1)

Data processing (ADD, SUB, CMP, MOV)

cond	Op	Rn	Rd	Shift	Rm	Rd := Rn Op Rm*
1				S		Rd := Rn Op S

ALU operation

Rm* = Rm with optional shift

multiply instructions are special case

Data transfer (to or from memory LDR, STR)

cond	Trans	Rn	Rd	S	Rd ↔ mem[Rn+/-S]
1				Shift Rm	Rd ↔ mem[Rn+/-Rm*]

Byte/word, load/store, etc

Multiple register transfer

cond	Type	Rn	Register list	Transfer registers to/from stack
------	------	----	---------------	----------------------------------

Overview (2)

Branch B, BL, BNE, BMI...

cond	L	S	PC := PC+8+4*S
0			L = 0 => Branch, B ...
1			L = 1 => Branch and link (R14 := PC+4), BL ...

S is sign extended
 NB +8 because of
 pipelining.

cond	coprocessor interface
1 1 0 0	
1 1 0 1	
1 1 1 0	

Software Interrupt (SWI)

cond	S	Simulate hardware interrupt: S is passed to handler in swi mode
1 1 1 1		

Paper Number(s): **EE1-9B**

IMPERIAL COLLEGE LONDON

DEPARTMENT OF ELECTRICAL AND ELECTRONIC ENGINEERING
EXAMINATIONS 2011

ISE Part I: MEng, BEng and ACGI

Corrected Copy

INTRODUCTION TO COMPUTER ARCHITECTURE AND SYSTEMS (PART B)
OPERATING SYSTEMS

Monday, 6 June 3.30 pm

Time allowed: 1:00 hour

There is ONE question on this paper, which must be answered.

Any special instructions for invigilators and information for candidates are on page 1.

Examiners responsible:

First Marker(s): Demir, Y.K.

Second Marker(s): Bouganis, C.

There is only ONE question below
Answer ALL parts

- (a) Consider the following set of processes, with their corresponding arrival times, duration, and priority levels [*higher numbers indicate higher priority*]:

Process	Arrival time (ms)	Duration (ms)	Priority level
A	0	2	1
B	2	6	2
C	5	2	3
D	7	4	4
E	9	2	5

Show the order of execution (including timing information) of the processes if the scheduler implements the following scheduling algorithms:

- i. Round Robin with a time slice of 4 ms [2]
- ii. Shortest Remaining Job First (SRJF) [2]
- iii. Priority-scheduling with pre-emption [2]

For each of the algorithms calculate the average waiting time, and the average turnaround time.

- (b) In the context of a memory paging system, consider the following scenario:

- You have three available frames
- The reference string is 3-4-1-8-1-2-3-6-5-4

Starting with empty frame contents, show the sequence of frame contents after each request, and count the number of page faults for each of the following page replacement algorithms:

- i. Optimal-page replacement [3]
- ii. First in First Out replacement [3]
- iii. LRU (Least Recently Used) page replacement [3]

- (c) A system has 23 instances of a resource type and there are currently four processes running; their maximum needs and their current allocation are shown in the table below

Process	Maximum requirements	Current allocation
A	10	5
B	13	7
C	5	2
D	10	4
Free: 5		

- i. Determine whether the current state is a safe state, or not, and in either case, demonstrate why. [2]
- ii. Describe the banker's algorithm for dynamic deadlock avoidance, and explain its main weakness. Provide the algorithm's response if process D requests three instances of the resource type. [3]

Introduction to Computer Architecture - Answers 2011

All questions are compulsory, marks are 40/30/30 (or 20/15/15 on whole paper) for Q1/Q2/Q3.

Questions will be slightly easier than in previous papers, due to lack of choice.

Answer to Question 1

Q1 is an easy question testing basic knowledge & understanding.

1.

a)

(i) &C1FA0000

(ii) $N = 8 \Rightarrow -1$, $N = 16 \Rightarrow 255$

[8]

b)

size = $16 * 4 * 512 = 32768$ bytes

select = $\log_2(16 * 4) = 6$ (includes LS 2 bits byte select) A5:0

index = 9 bits: A14:6

tag = A31:15

[8]

c)

```
SUB R3, R1, R3
ADD R3, R3, R2
ADD R3, R3, #111
```

Many other equivalents. Note 1st line must include R3

[8]

d)

(i) A (R0)

(ii) C

(iii) B (R0)

(iv) B

[8]

(e)

ADD R10, R10, R9, lsl #12

ADD R3, R3, R3, lsr #3

[8]

Answer to Question 2

This question tests ability to understand and analyse operation of ARM assembly code in detail. It requires accuracy and comprehensive understanding of the instructions, but is straightforward.

- (a) tests understanding of two's complement arithmetic, and condition codes
- (b) test understanding of memory addressing modes & logical operations
- (c) tests understanding of shift, byte transfer, addressing modes & simple loops

Location	Value
&100	&04030201
&104	&08070605
> &10C	&0

Figure 2.1. Memory locations

```
MOV R0, #1
MVN R1, R0
ADDS R2, R0, R1
SBC R3, R0, R1
ADC R0, R0, R0
```

(a)

```
MOV R10, &100
LDR R0, [R10], #4
LDR R1, [R10], #4
ORR R2, R0, R1
ANDS R3, R0, R1
```

(b)

```
MOV R2, #3
MOV R1, #&100
L1 LDRB R0, [R1, #1]!
ORR R3, R0, R3, LSL #8
SUBS R2, R2, #1
BGE L1
```

(c)

Figure 2.2. Code fragments

2 marks each box

	R0	R1	R2	R3	NZCV
(a)	2	FFFFFFFE	FFFFFFFF	2	1000
(b)	04030201	08070605	08070605	00030201	0000
(c)	8	104	FFFFFFFF	03020108	1000

Figure 2.3. Template for answers

Answer to Question 3

This more difficult question tests knowledge of the ARM subroutine implementation, as well as instruction timing, and ability to manipulate code.

a)

STM/LDM instructions save/restore data on a downwards growing empty stack pointer stack. R13 is the stack pointer. In this case R2,R14 are saved, R2,R15 are restored - the restore operation thus causes the subroutine return by loading PC with the return address previously in R14.

[6]

b)

(i) $4+5+1+1+1+1+1+(5+3)+4 = 26$ cycles

(ii) 24 cycles

STMED takes longer because:

- It is a random access memory operation which slows down the memory unit (precise reason here not needed).
- It has destination PC which breaks the prefetch/decode pipeline.

[8]

c)

R2(0) has values: R1(0)->xor R1(1:0)->xor R1(3:0) -> xor R1(7:0)

Z=1 iff R2(0)=0 => xor of R1(7:0) is 0.

[8]

d)

PARITY

```

; note MOV & LDM are not needed
A      EOR    R1, R1, R1, ror #1
        EOR    R1, R1, R1, ror #2
        EOR    R1, R1, R1, ror #4
B      ANDS   R1, R1, #1
        MOV    R15, R14 ; note MOV not LDM

START  MOV    R2,R1
        BL     PARITY
        ORRNE  R1, R2, #&100 ; note optimisation
FINISH
```

[8]

Question 1(a) *New computed example*

[R= Running, - = waiting]

Round Robin – 4 ms

A	R	R																	
B			R	R	R	R	-	-	R	R									
C						-	R	R											
D								-	-	-	R	R	R	R					
E									-	-	-	-	-	-	R	R			

Avg waiting time: $(2+1+3+5) / 5 = 2.2$ ms, Avg turnaround time: $27/5 = 5.4$ ms

[2]

SRJF

A	R	R																	
B			R	R	R	-	-	R	R	R									
C						R	R												
D								-	-	-	-	-	R	R	R	R			
E									-	R	R								

AWT = $2+5+1/5 = 1.6$ ms ATT = $24/5 = 4.8$ ms

[2]

Priority Scheduling (with preemption)

A	R	R																	
B			R	R	R	-	-	-	-	-	-	-	R	R	R				
C						R	R												
D								R	R	-	-	R	R						
E									R	R									

AWT = $10/5 = 2$ ms; ATT = $26/5 = 5.2$ ms

[2]

(b) [new computed example]

Optimal page replacement algorithm (5 page faults)

	3	4	3	8	1	2	3	3	4	2
Frame1	3	3		3	3	3				
Frame2	-	4		4	4	4				
Frame3	-	-		8	1	2				

[3]

FIFO page replacement algorithm (8 page faults)

	3	4	3	8	1	2	3	3	4	2
Frame1	3	3		3	1	1	3		3	3
Frame2	-	4		4	4	2	2		4	5
Frame3	-	-		8	8	8	8		8	2

[3]

LRU (Least recently used) page replacement algorithm (7 page faults)

	3	4	3	8	1	2	3	3	4	2
Frame1	3	3		3	3	2	2		2	
Frame2	-	4		4	1	1	1		4	
Frame3	-	-		8	8	8	3		3	

[3]

(c) [New computed example]

i. Current state is a safe state since there is a safe sequence A->B->C->D that can be followed when allocating resources to processes, in order for a deadlock not to be possible. [2]

ii. Banker's algorithm works as follows: when a process requests a resource, the algorithm first determines whether granting the request will lead to an unsafe state – if doesn't it grants the request, otherwise the decision is postponed until a process releases some of its resources. To check whether the state is safe, the algorithm:

- checks whether it has some resources to satisfy some process
- The resources of that process are presumed released and added to the available resources
- steps 1 and 2 are repeated until we find that all current processes can be satisfied.

The algorithm's main weakness is that it needs to know the resource requirements for each process in advance.

In this particular case, a request by process D for 3 resources will be denied given that the system will subsequently be in an unsafe state, since there is no safe sequence of satisfying processes with the appropriate resources. [3]