

Computer Architecture EE2-13

Department of Computing Examinations — 2015–2016 Session

Confidential

MODEL ANSWER and MARKING SCHEME

Examiner

wl

Paper Code

C210

Question 1

Page 1

out of 1

Question labels in left margin

Mark allocations in right margin

a $m \times 9 = 8m + m$, so $\text{sll } \$2, \$1, 3$
 $\text{add } \$2, \$1, \$2$
 Take 3 cycles, $5/3 = 1.67$ times faster than the 'mult' instruction 6

b If $k = 1022 = \underbrace{11111111}_9 0$, need 9 shifts and 8 adds
 Let $\$1$ contains the number m , $\$2$ contains the result $9m$
 $\$3$ provides temporary storage

```

sll $2 $1 1
sll $3 $1 2
1 add $2 $2 $3
  sll $3 $1 3
2 add $2 $2 $3
  sll $3 $1 4
3 add $2 $2 $3
  sll $3 $1 5
4 add $2 $2 $3
  sll $3 $1 6
5 add $2 $2 $3
  sll $3 $1 7
6 add $2 $2 $3
  sll $3 $1 8
7 add $2 $2 $3
  sll $3 $1 9
8 add $2 $2 $3
  
```

9 shifts and 8 adds

This takes $9n + 8(2n) = 25n$,

5 times slower than 'mult' instruction 6

c Using Booth's algorithm, replace the 8 adds by 1 subtract, since
 $m \times 1022 = m(1024 - 2) = 1024m - 2m$

multiply by 1024 is shift left by 10 bits, so

```

sll $2 $1 2
sll $3 $1 10
sub $2 $2 $2
  
```

} This takes 4n cycles, $\frac{5}{4}$ times
 faster than the mult instruction 8

Department of Computing Examinations – 2015 - 2016 Session		Confidential
MODEL ANSWERS and MARKING SCHEME		
First Examiner: David Thomas		Second Examiner: Wayne Luk
Paper: C210=E2.13 - Computer Architecture	Question: 2	Page 1 of 3

- 2a i) Considering L1 vs L3 cache levels, describe two properties of caches and how you would expect them to vary between L1 and L3.

- * The capacity of the L1 cache will be smaller than the L3.
- * The hit time of the L1 cache will be lower than the L3.
- * The associativity of the L1 cache will be lower than the L3.

Marks:

2

- ii) Give a fragment of code that demonstrates both temporal locality and spatial locality. You can use any language as long as semantics are clear.

```
uint32_t *pX = ...;
uint32_t acc=0;
for(int i=0;i<100;i++){
    // Loops over sequential words (spatial)
    // Visits each word 10 times (temporal)
    acc=acc+pX[i%10];
}
```

Marks:

3

- iii) Assume a 2-way set associative cache with w bytes per word, b words per block, and a total cache size of c bytes. Given a byte_address, give statements for calculating the set_index. Use intermediate variables that show your working.

```
word_address = byte_address / w;
block_address = word_address / b;
num_sets = c / (w*b*2);
set_index = block_address % num_sets;
```

Marks:

4

- b A CPU designer has suggested that the MIPS ISA is extended with an integer multiply-accumulate instruction, with the following syntax and semantics:

```
mul_add $a, $b, $c // a = a + b * c
```

where \$a, \$b, and \$c are all registers.

- i) How many data operand registers does a MIPS R-type instruction support?

Three (1 destination; 2 source).

Marks:

2

Department of Computing Examinations – 2015 - 2016 Session		Confidential
MODEL ANSWERS and MARKING SCHEME		
First Examiner: David Thomas	Second Examiner: Wayne Luk	
Paper: C210=E2.13 - Computer Architecture	Question: 2	Page 2 of 3

- ii) How would the register-file of a 5-stage pipelined MIPS need to be modified in order to support this instruction?

At the moment the register-file only supports the reading of two operands per cycle. This instruction needs to read three operands, so another read port would need to be added.

Marks:

2

- iii) Give a sequence of standard MIPS instructions that is equivalent to `mul_add`. You may use the register `$t` as a temporary register.

```
mulu $b, $c
nop          // Needed according to MIPS ISA
mflo $t
add $a, $a, $t
```

Marks:

3

- iv) Assume that both the standard MIPS and extended MIPS execute one instruction per cycle, and caches are warm. Provide a quantitative estimate for the speedup that `mul_add` could provide for the following function. State any assumptions, with a brief explanation of why they are reasonable.

```
uint32_t dot_product(uint32_t n, uint32_t *a, uint32_t *b)
{
    uint32_t acc=0;
    for(uint32_t i=0; i<n; i++){
        acc = acc + a[i] * b[i];
    }
    return acc;
}
```

The inner loop of the function must contain:

- * Two loads
- * A multiply accumulate
- * Two loop increments
- * A branch.

My assumptions (others could be acceptable, e.g. loop pipelining) are:

Department of Computing Examinations – 2015 - 2016 Session		Confidential
MODEL ANSWERS and MARKING SCHEME		
First Examiner: David Thomas		Second Examiner: Wayne Luk
Paper: C210=E2.13 - Computer Architecture	Question: 2	Page 3 of 3

- * The compiler can avoid load-use stalls, for example by loop skewing.
- * No loop pipelining is used, as the compiler is optimising for code size.
- * One instruction can be placed in the delay slot, as there are a couple of increments.
- * For the standard MIPS, I assume one instruction can be placed in the NOP.

This results in standard MIPS = $2+3+2+1=8$, and extended MIPS = $2+1+2+1=6$.

So asymptotically an estimate of speed-up is $8/6 = 4/3 = 1.25x$.

Marks:

4

The two parts carry, respectively, 45%, and 55% of the marks.