

IMPERIAL COLLEGE LONDON

DEPARTMENT OF ELECTRICAL AND ELECTRONIC ENGINEERING
EXAMINATIONS 2009

EEE/ISE PART III/IV: MEng, BEng and ACGI

REAL-TIME OPERATING SYSTEMS

Friday, 15 May 10:00 am

Time allowed: 3:00 hours

There are SIX questions on this paper.

Answer FOUR questions.

All questions carry equal marks.

Any special instructions for invigilators and information for candidates are on page 1.

Examiners responsible	First Marker(s) :	T.J.W. Clarke
	Second Marker(s) :	Y.K. Demiris

Special instructions for invigilators

The booklet RTOS Exam Notes 2009 should be distributed with the paper.

Special instructions for students

You may use the booklet RTOS Exam Notes 2009 which is a reproduction of that published on the course web-site before the exam.

The Questions

1.

- (a) Explain carefully the different characteristics of deadlock, starvation, and livelock in a real-time system. [8]
- (b) Write, using pseudo-code, an example application which might suffer from real livelock. Suggest, with modified pseudo-code, how the livelock in your example can be detected and prevented. [8]
- (c) An application uses busy-wait loops to ensure exclusive resource use. It is observed that all tasks continue to run, but no progress is made. Comment on possible reasons for this. [4]

2. This question relates to the v4.0.5 FreeRTOS implementation of queues: source code for FreeRTOS v4.0.5 is contained in the Exam Notes 2009.

- (a) Explain with an appropriate timeline what is the sequence of task states of task X sending two messages, and task Y receiving 1 message, communicating through a message queue of length 1 which is initially empty, in such a way that both tasks X and Y block temporarily. State on your diagram what causes each state transition. [5]
- (b) Identify the source code lines in Exam Notes 2009 in the execution *QueueSend()* and *QueueReceive()* at each point in your timeline. [5]
- (c) State under precisely what circumstances the FreeRTOS v4.0.5 code could result in a call to *QueueReceive()* returning with an error code when there has been no timeout. [5]
- (d) Considering the case of one sender and multiple receivers on an empty queue discuss what semantics for *QueueReceive()* would be ideal in a pre-emptive priority-based system, where higher priority tasks should if possible have priority in message queues. Discuss why implementation of this would be difficult. [5]

3.

(a) A priority-scheduled real-time system consists of four jobs with the characteristics shown in *Figure 3.1*. Answer the following questions, using inequalities on $t1$, $t2$ as necessary to determine the answer. You may assume RTOS task-switching overheads are negligible.

(i) How would you prioritise these jobs under an RTOS with prioritised tasks? (Your answer may depend on $t1$ and $t2$).

(ii) Assuming the prioritisation in (i), state for which values of $t1$, $t2$ the following cases apply:

- A. The system will certainly not meet all deadlines
- B. It is not known whether the system will meet all deadlines
- C. The system is guaranteed to meet all deadlines.

[8]

(b) Suppose that the jobs W, Z share mutually exclusive access to a resource R. The resource is claimed for a 20 us critical section once in each job period by both W & Z. Determine what are the appropriate blocking times to use in extended RMA for this system, and hence establish an inequality on $t1, t2$ which guarantees all deadlines are met in the new system.

[6]

(c) Suppose that in the system described by *Figure 3.1* and part (a) job Z has higher priority than W, all other relative priorities are as in your answer to part (a) (i). In the revised system calculate an inequality on $t1, t2$ which guarantees all deadlines by finding an upper limit on the time that W is unable to execute because Z is executing and treating this as blocking time.

[6]

Job	Job Time	Job Period
X	50 us	$t1$ us
Y	100 us	$t2$ us
Z	50 us	500 us
W	20 us	200 us

Figure 3.1

4.

Figure 4.1 shows part of the code for tasks A,B,C & D in a FreeRTOS application. There are no calls to *printer_open()*, *printer_close()* and *print()* except as indicated in Figure 4.1. It is specified that:

- (1) A,B,C have all executed *printer_open()* more recently than *printer_close()* whenever function *print()* is called.
- (2) A,B,C have all executed *printer_close()* after *print()* returns before any one of them executes *printer_open()* or *print()* is called again.
- (3) *printer_open()* and *printer_close()* are never active at the same time as *print()*.

Write FreeRTOS code for segments S_{A1} , S_{B1} , S_{C1} , S_{A2} , S_{B2} , S_{C2} and S_{D1} , S_{D2} , and a function *initialise()* called once before any of the code in Figure 4.1 is executed which ensures that conditions (1), (2) and (3) above always ahold. Explain clearly why your code ensures each of these three conditions is satisfied.

[20]

Task A	Task B	Task C	Task D
....
<i>printer_open()</i>	<i>printer_open()</i>	<i>printer_open()</i>	S_{D1}
S_{A1}	S_{B1}	S_{C1}	<i>print()</i>
<i>printer_close()</i>	<i>printer_close()</i>	<i>printer_close()</i>	S_{D2}
S_{A2}	S_{B2}	S_{C2}
.....	

Figure 4.1

5. A game is played with N rods labelled $1, 2, \dots, N$. Rod x can contain any non-negative number of beads $B(x)$. A single move consists of examining the number of beads on two rods, x and y , and changing the number of beads on x, y simultaneously to new values according to a function f :

$$B(x) := f(B(x), B(y))$$

$$B(y) := f(B(y), B(x))$$

The two calls to function $f()$ are both made using the old values of $B(x), B(y)$, and the function $f()$ has a return value which depends only on its parameters.

It is proposed to implement a timed simulation of this strategy with a number of tasks each implementing a sequence of moves using randomly chosen rods, and locations in shared memory implementing the values $B(x)$. Multiple moves can be made in parallel by different tasks as long as each move uses a disjoint set of rods.

- (a) Using scheduler locking write pseudo-code for a function:

Move(int x, int y)

which implements a single move on rods x, y and ensures both that the memory values $B(i)$ are always consistent and that moves which interfere with each other will therefore never be made.

[4]

- (b) Using semaphores write an implementation of **Move(int x, int y)**. State in what ways this is better, and in what ways worse, than your answer to (a).

[10]

- (c) Prove that your answer to part (b) cannot deadlock.

[6]

6.

- (a) Explain the difference between scheduler locking and interrupt locking in FreeRTOS. [4]
- (b) State three ways in which a critical section can be implemented in an RTOS. In each case determine under what circumstances the implementation will increase overall worst-case task or interrupt latency. [6]
- (c) Explain precisely how a clock tick is implemented in FreeRTOS 4.0.5, illustrating the sequence of operations with an appropriate diagram, if a clock tick interrupt occurs:
 - (i) while interrupts are locked.
 - (ii) while the scheduler is locked.Detail changes to any data structures affected by the transaction. [10]

[END]

RTOS EXAM NOTES 2009

Task.h

```

1  typedef void * xTaskHandle;
2
3  #define taskYIELD()          portYIELD()
4  #define taskENTER_CRITICAL() portENTER_CRITICAL()
5  #define taskEXIT_CRITICAL()  portEXIT_CRITICAL()
6  #define taskDISABLE_INTERRUPTS() portDISABLE_INTERRUPTS()
7  #define taskENABLE_INTERRUPTS() portENABLE_INTERRUPTS()
8
9
10 /*-----
11  * TASK CREATION API
12  *-----*/
13
14 signed portBASE_TYPE xTaskCreate(    pdTASK_CODE pvTaskCode, const signed portCHAR * const pcName,
15                                     unsigned portSHORT usStackDepth, void *pvParameters,
16                                     unsigned portBASE_TYPE uxPriority, xTaskHandle *pvCreatedTask );
17
18 void vTaskDelete( xTaskHandle pxTask );
19
20 /*-----
21  * TASK CONTROL API
22  *-----*/
23
24 void vTaskDelay( portTickType xTicksToDelay );
25 void vTaskDelayUntil( portTickType *pxPreviousWakeTime, portTickType xTimeIncrement );
26 unsigned portBASE_TYPE uxTaskPriorityGet( xTaskHandle pxTask );
27 void vTaskPrioritySet( xTaskHandle pxTask, unsigned portBASE_TYPE uxNewPriority );
28 void vTaskSuspend( xTaskHandle pxTaskToSuspend );
29 void vTaskResume( xTaskHandle pxTaskToResume );
30 portBASE_TYPE xTaskResumeFromISR( xTaskHandle pxTaskToResume );
31
32 /*-----
33  * SCHEDULER CONTROL
34  *-----*/
35
36 void vTaskStartScheduler( void );
37 void vTaskEndScheduler( void );
38 void vTaskSuspendAll( void );
39 signed portBASE_TYPE xTaskResumeAll( void );
40
41 /*-----
42  * TASK UTILITIES
43  *-----*/
44
45 portTickType xTaskGetTickCount( void );
46 unsigned portBASE_TYPE uxTaskGetNumberOfTasks( void );
47 void vTaskPlaceOnEventList( xList *pxEventList, portTickType xTicksToWait );
48 signed portBASE_TYPE xTaskRemoveFromEventList( const xList *pxEventList );
49 void vTaskCleanUpResources( void );
50 inline void vTaskSwitchContext( void );
51 xTaskHandle xTaskGetCurrentTaskHandle( void );
52
53
54

```

Semaphr.c

```

56
57 #define vSemaphoreCreateBinary( xSemaphore ) { \
58     xSemaphore = xQueueCreate( ( unsigned portCHAR ) 1, semSEMAPHORE_QUEUE_ITEM_LENGTH ); \
59     if( xSemaphore != NULL ) \
60     { \
61         xSemaphoreGive( xSemaphore ); \
62     } \
63 }
64
65 #define xSemaphoreTake( xSemaphore, xBlockTime ) \
66     xQueueReceive( ( xQueueHandle ) xSemaphore, NULL, xBlockTime )
67
68 #define xSemaphoreGive( xSemaphore ) xQueueSend( ( xQueueHandle ) xSemaphore, NULL, semGIVE_BLOCK_TIME )
69
70 #define xSemaphoreGiveFromISR( xSemaphore, xTaskPreviouslyWoken ) \
71     xQueueSendFromISR( ( xQueueHandle ) xSemaphore, NULL, xTaskPreviouslyWoken )

```

From Task.h - related to lists package

```

100 signed portBASE_TYPE xTaskRemoveFromEventList( const xList *pxEventList )
101 {
102     tskTCB *pxUnblockedTCB;
103     portBASE_TYPE xReturn;
104
105     /* THIS FUNCTION MUST BE CALLED WITH INTERRUPTS DISABLED OR THE
106     SCHEDULER SUSPENDED. It can also be called from within an ISR. */
107
108     /* The event list is sorted in priority order, so we can remove the
109     first in the list, remove the TCB from the delayed list, and add
110     it to the ready list.
111
112     If an event is for a queue that is locked then this function will never
113     get called - the lock count on the queue will get modified instead. This
114     means we can always expect exclusive access to the event list here. */
115     pxUnblockedTCB = ( tskTCB * ) listGET_OWNER_OF_HEAD_ENTRY( pxEventList );
116     vListRemove( &( pxUnblockedTCB->xEventListItem ) );
117
118     if( uxSchedulerSuspended == ( unsigned portBASE_TYPE ) pdFALSE )
119     {
120         vListRemove( &( pxUnblockedTCB->xGenericListItem ) );
121         prvAddTaskToReadyQueue( pxUnblockedTCB );
122     }
123     else
124     {
125         /* We cannot access the delayed or ready lists, so will hold this
126         task pending until the scheduler is resumed. */
127         vListInsertEnd( ( xList * ) &( xPendingReadyList ), &( pxUnblockedTCB->xEventListItem ) );
128     }
129
130     if( pxUnblockedTCB->uxPriority >= pxCurrentTCB->uxPriority )
131     {
132         /* Return true if the task removed from the event list has
133         a higher priority than the calling task. This allows
134         the calling task to know if it should force a context
135         switch now. */
136         xReturn = pdTRUE;
137     }
138     else
139     {
140         xReturn = pdFALSE;
141     }
142
143     return xReturn;
144 }

```

```

147 List.h
148 /*
149  * Definition of the only type of object that a list can contain.
150  */
151 struct xLIST_ITEM
152 {
153     portTickType xItemValue; /*< The value being listed. In most cases this is
154                               used to sort the list in descending order. */
155     volatile struct xLIST_ITEM * pxNext; /*< Pointer to the next xListItem in the list. */
156     volatile struct xLIST_ITEM * pxPrevious; /*< Pointer to the previous xListItem in the list. */
157     void * pvOwner; /*< Pointer to the object (normally a TCB) that contains the list item. */
158     void * pvContainer; /*< Pointer to the list in which this list item is placed (if any). */
159 };
160 typedef struct xLIST_ITEM xListItem; /* For some reason lint wants this as two separate definitions. */
161
162 struct xMINI_LIST_ITEM
163 {
164     portTickType xItemValue;
165     volatile struct xLIST_ITEM * pxNext;
166     volatile struct xLIST_ITEM * pxPrevious;
167 };
168 typedef struct xMINI_LIST_ITEM xMiniListItem;
169
170 /*
171  * Definition of the type of queue used by the scheduler.
172  */
173 typedef struct xLIST
174 {
175     volatile unsigned portBASE_TYPE uxNumberOfItems;
176     volatile xListItem * pxIndex; /* Used to walk through the list */
177     volatile xMiniListItem xListEnd; /* List item that contains the maximum possible item value */
178 } xList;
179
180 #define listSET_LIST_ITEM_OWNER( pxListItem, pxOwner ) ( pxListItem )->pvOwner = ( void * ) pxOwner
181
182 #define listSET_LIST_ITEM_VALUE( pxListItem, xValue ) ( pxListItem )->xItemValue = xValue
183
184 #define listGET_LIST_ITEM_VALUE( pxListItem ) ( ( pxListItem )->xItemValue )
185
186 #define listLIST_IS_EMPTY( pxList ) ( ( pxList )->uxNumberOfItems == ( unsigned portBASE_TYPE ) 0 )
187
188 #define listCURRENT_LIST_LENGTH( pxList ) ( ( pxList )->uxNumberOfItems )
189
190 #define listGET_OWNER_OF_NEXT_ENTRY( pxTCB, pxList ) \
191     /* Increment the index to the next item and return the item, ensuring */ \
192     /* we don't return the marker used at the end of the list. */ \
193     ( pxList )->pxIndex = ( pxList )->pxIndex->pxNext; \
194     if( ( pxList )->pxIndex == ( xListItem * ) &( ( pxList )->xListEnd ) ) \
195     { \
196         ( pxList )->pxIndex = ( pxList )->pxIndex->pxNext; \
197     } \
198     pxTCB = ( pxList )->pxIndex->pvOwner
199
200
201 #define listGET_OWNER_OF_HEAD_ENTRY( pxList ) ( ( pxList->uxNumberOfItems != ( unsigned portBASE_TYPE ) 0 \
202 ) ? ( ( &( pxList->xListEnd ) )->pxNext->pvOwner ) : ( NULL ) )
203
204 #define listIS_CONTAINED_WITHIN( pxList, pxListItem ) ( ( pxListItem )->pvContainer == ( void * ) pxList )
205
206 void vListInitialise( xList *pxList );
207
208 void vListInitialiseItem( xListItem *pxItem );
209
210 void vListInsert( xList *pxList, xListItem *pxNewListItem );
211
212 void vListInsertEnd( xList *pxList, xListItem *pxNewListItem );
213
214 void vListRemove( xListItem *pxItemToRemove );
215

```

```

216 List.c
217 #include <stdlib.h>
218 #include "FreeRTOS.h"
219 #include "list.h"
220
221 /*-----
222  * PUBLIC LIST API documented in list.h
223  *-----*/
224
225 void vListInitialise( xList *pxList )
226 {
227     /* The list structure contains a list item which is used to mark the end of the list. To initialise
228        the list the list end is inserted as the only list entry. */
229     pxList->pxIndex = ( xListItem * ) &(amp; pxList->xListEnd );
230
231     /* The list end value is the highest possible value in the list to ensure it
232        remains at the end of the list. */
233     pxList->xListEnd.xItemValue = portMAX_DELAY;
234
235     /* The list end next and previous pointers point to itself so we know when the list is empty. */
236     pxList->xListEnd.pxNext = ( xListItem * ) &(amp; pxList->xListEnd );
237     pxList->xListEnd.pxPrevious = ( xListItem * ) &(amp; pxList->xListEnd );
238
239     pxList->uxNumberOfItems = 0;
240 }
241
242 void vListInitialiseItem( xListItem *pxItem )
243 {
244     /* Make sure the list item is not recorded as being on a list. */
245     pxItem->pvContainer = NULL;
246 }
247
248 void vListInsertEnd( xList *pxList, xListItem *pxNewListItem )
249 {
250     volatile xListItem * pxIndex;
251
252     /* Insert a new list item into pxList, but rather than sort the list, makes the new list item the
253        last
254        item to be removed by a call to pvListGetOwnerOfNextEntry. This means it has to be the item
255        pointed to by the pxIndex member. */
256     pxIndex = pxList->pxIndex;
257
258     pxNewListItem->pxNext = pxIndex->pxNext;
259     pxNewListItem->pxPrevious = pxList->pxIndex;
260     pxIndex->pxNext->pxPrevious = ( volatile xListItem * ) pxNewListItem;
261     pxIndex->pxNext = ( volatile xListItem * ) pxNewListItem;
262     pxList->pxIndex = ( volatile xListItem * ) pxNewListItem;
263
264     /* Remember which list the item is in. */
265     pxNewListItem->pvContainer = ( void * ) pxList;
266
267     ( pxList->uxNumberOfItems )++;
268 }
269
270

```

```

271 void vListInsert( xList *pxList, xListItem *pxNewListItem )
272 {
273     volatile xListItem *pxIterator;
274     portTickType xValueOfInsertion;
275
276     /* Insert the new list item into the list, sorted in ulListItem order. */
277     xValueOfInsertion = pxNewListItem->xItemValue;
278
279     /* If the list already contains a list item with the same item value then the new list item should be
280     placed after it. This ensures that TCB's which are stored in ready lists (all of which have the same
281     ulListItem value) get an equal share of the CPU. However, if the xItemValue is the same as the back
282     marker the iteration loop below will not end. This means we need to guard against this by checking
283     the value first and modifying the algorithm slightly if necessary. */
284     if( xValueOfInsertion == portMAX_DELAY )
285     {
286         pxIterator = pxList->xListEnd.pxPrevious;
287     }
288     else
289     {
290         for( pxIterator = ( xListItem * ) &(amp; pxList->xListEnd );
291             pxIterator->pxNext->xItemValue <= xValueOfInsertion;
292             pxIterator = pxIterator->pxNext )
293         {
294             /* There is nothing to do here, we are just iterating to the wanted insertion position. */
295         }
296     }
297
298     pxNewListItem->pxNext = pxIterator->pxNext;
299     pxNewListItem->pxNext->pxPrevious = ( volatile xListItem * ) pxNewListItem;
300     pxNewListItem->pxPrevious = pxIterator;
301     pxIterator->pxNext = ( volatile xListItem * ) pxNewListItem;
302
303     /* Remember which list the item is in. This allows fast removal of the item later. */
304     pxNewListItem->pvContainer = ( void * ) pxList;
305
306     ( pxList->uxNumberOfItems )++;
307 }
308
309 void vListRemove( xListItem *pxItemToRemove )
310 {
311     xList * pxList;
312     pxItemToRemove->pxNext->pxPrevious = pxItemToRemove->pxPrevious;
313     pxItemToRemove->pxPrevious->pxNext = pxItemToRemove->pxNext;
314
315     /* The list item knows which list it is in. Obtain the list from the list item. */
316     pxList = ( xList * ) pxItemToRemove->pvContainer;
317
318     /* Make sure the index is left pointing to a valid item. */
319     if( pxList->pxIndex == pxItemToRemove )
320     {
321         pxList->pxIndex = pxItemToRemove->pxPrevious;
322     }
323
324     pxItemToRemove->pvContainer = NULL;
325     ( pxList->uxNumberOfItems )--;
326 }
327 /*-----*/

```

Queue.h

```

typedef void * xQueueHandle;

xQueueHandle xQueueCreate( unsigned portBASE_TYPE uxQueueLength, unsigned portBASE_TYPE uxItemSize );

signed portBASE_TYPE xQueueSend( xQueueHandle xQueue, const void * pvItemToQueue, portTickType xTicksToWait );

signed portBASE_TYPE xQueueReceive( xQueueHandle xQueue, void *pvBuffer, portTickType xTicksToWait );

unsigned portBASE_TYPE uxQueueMessagesWaiting( xQueueHandle xQueue );

void vQueueDelete( xQueueHandle xQueue );

signed portBASE_TYPE xQueueSendFromISR( xQueueHandle pxQueue, const void *pvItemToQueue, signed portBASE_TYPE
xTaskPreviouslyWoken );

signed portBASE_TYPE xQueueReceiveFromISR( xQueueHandle pxQueue, void *pvBuffer, signed portBASE_TYPE
*pxTaskWoken );

```

Queue.c

```

/*-----
 * PUBLIC LIST API documented in list.h
 *-----*/

/* Constants used with the cRxLock and cTxLock structure members. */
#define queueUNLOCKED ( ( signed portBASE_TYPE ) -1 )

/*
 * Definition of the queue used by the scheduler.
 * Items are queued by copy, not reference.
 */
typedef struct QueueDefinition
{
    signed portCHAR *pcHead; /*< Points to the beginning of the queue storage area. */
    signed portCHAR *pcTail; /*< Points to the byte at the end of the queue storage area.
                               Once more byte is allocated than necessary to store the queue items,
                               this is used as a marker. */

    signed portCHAR *pcWriteTo; /*< Points to the free next place in the storage area. */
    signed portCHAR *pcReadFrom; /*< Points to the last place that a queued item was read from. */

    xList xTasksWaitingToSend; /*< List of tasks that are blocked waiting to post onto this queue.
                               Stored in priority order. */
    xList xTasksWaitingToReceive; /*< List of tasks that are blocked waiting to
                               read from this queue. Stored in priority order. */

    unsigned portBASE_TYPE uxMessagesWaiting; /*< The number of items currently in the queue. */
    unsigned portBASE_TYPE uxLength; /*< The length of the queue defined as the number
                                       of items it will hold, not the number of bytes. */
    unsigned portBASE_TYPE uxItemSize; /*< The size of each items that the queue will hold. */

    signed portBASE_TYPE xRxLock; /*< Stores the number of items received from the queue
                                   (removed from the queue) while the queue was locked.
                                   Set to queueUNLOCKED when the queue is not locked. */
    signed portBASE_TYPE xTxLock; /*< Stores the number of items transmitted to the queue
                                   (added to the queue) while the queue was locked.
                                   Set to queueUNLOCKED when the queue is not locked. */
} xQUEUE;

/*-----*/

/*
 * Inside this file xQueueHandle is a pointer to a xQUEUE structure.
 * To keep the definition private the API header file defines it as a
 * pointer to void.
 */
typedef xQUEUE * xQueueHandle;

```

```

369 /*
370 * Unlocks a queue locked by a call to prvLockQueue. Locking a queue does not
371 * prevent an ISR from adding or removing items to the queue, but does prevent
372 * an ISR from removing tasks from the queue event lists. If an ISR finds a
373 * queue is locked it will instead increment the appropriate queue lock count
374 * to indicate that a task may require unblocking. When the queue is unlocked
375 * these lock counts are inspected, and the appropriate action taken.
376 */
377 static signed portBASE_TYPE prvUnlockQueue( xQueueHandle pxQueue );
378
379 /*
380 * Uses a critical section to determine if there is any data in a queue.
381 *
382 * @return pdTRUE if the queue contains no items, otherwise pdFALSE.
383 */
384 static signed portBASE_TYPE prvIsQueueEmpty( const xQueueHandle pxQueue );
385
386 /*
387 * Uses a critical section to determine if there is any space in a queue.
388 *
389 * @return pdTRUE if there is no space, otherwise pdFALSE;
390 */
391 static signed portBASE_TYPE prvIsQueueFull( const xQueueHandle pxQueue );
392
393 /*
394 * Macro that copies an item into the queue. This is done by copying the item
395 * byte for byte, not by reference. Updates the queue state to ensure it's
396 * integrity after the copy.
397 */
398 #define prvCopyQueueData( pxQueue, pvItemToQueue ) \
399 { \
400     memcpy( ( void * ) pxQueue->pcWriteTo, pvItemToQueue, ( unsigned ) pxQueue->uxItemSize ); \
401     ++( pxQueue->uxMessagesWaiting ); \
402     pxQueue->pcWriteTo += pxQueue->uxItemSize; \
403     if( pxQueue->pcWriteTo >= pxQueue->pcTail ) \
404     { \
405         pxQueue->pcWriteTo = pxQueue->pcHead; \
406     } \
407 } \
408
409 /*
410 * Macro to mark a queue as locked. Locking a queue prevents an ISR from accessing the queue event lists.
411 */
412 #define prvLockQueue( pxQueue ) \
413 { \
414     taskENTER_CRITICAL(); \
415     ++( pxQueue->xRxLock ); \
416     ++( pxQueue->xTxLock ); \
417     taskEXIT_CRITICAL(); \
418 } \
419
420 /*-----
421 * PUBLIC QUEUE MANAGEMENT API documented in queue.h
422 *-----*/
423
424 xQueueHandle xQueueCreate( unsigned portBASE_TYPE uxQueueLength, unsigned portBASE_TYPE uxItemSize )
425 {
426     xQUEUE *pxNewQueue;
427     size_t xQueueSizeInBytes;
428
429     /* Allocate the new queue structure. */
430     if( uxQueueLength > ( unsigned portBASE_TYPE ) 0 )
431     {
432         pxNewQueue = ( xQUEUE * ) pvPortMalloc( sizeof( xQUEUE ) );
433         if( pxNewQueue != NULL )
434         {
435             /* Create the list of pointers to queue items. The queue is one byte
436              longer than asked for to make wrap checking easier/faster. */
437             xQueueSizeInBytes = ( size_t ) ( uxQueueLength * uxItemSize ) + ( size_t ) 1;
438
439             pxNewQueue->pcHead = ( signed portCHAR * ) pvPortMalloc( xQueueSizeInBytes );
440             if( pxNewQueue->pcHead != NULL )
441             {
442                 /* Initialise the queue members as described above where the
443                  queue type is defined. */
444                 pxNewQueue->pcTail = pxNewQueue->pcHead + ( uxQueueLength * uxItemSize );
445                 pxNewQueue->uxMessagesWaiting = 0;
446                 pxNewQueue->pcWriteTo = pxNewQueue->pcHead;
447                 pxNewQueue->pcReadFrom = pxNewQueue->pcHead + ( ( uxQueueLength - 1 ) *
448                                                                 uxItemSize );
449                 pxNewQueue->uxLength = uxQueueLength;
450                 pxNewQueue->uxItemSize = uxItemSize;
451                 pxNewQueue->xRxLock = queueUNLOCKED;
452                 pxNewQueue->xTxLock = queueUNLOCKED;

```

```

452
453         /* Likewise ensure the event queues start with the correct state. */
454         vListInitialise( &(amp; pxNewQueue->xTasksWaitingToSend ) );
455         vListInitialise( &(amp; pxNewQueue->xTasksWaitingToReceive ) );
456
457         return pxNewQueue;
458     }
459     else
460     {
461         vPortFree( pxNewQueue );
462     }
463 }
464
465 /* Will only reach here if we could not allocate enough memory or no memory
466 was required. */
467 return NULL;
468 }
469
470 signed portBASE_TYPE xQueueSend( xQueueHandle pxQueue, const void *pvItemToQueue, portTickType xTicksToWait )
471 {
472     signed portBASE_TYPE xReturn;
473
474     /* Make sure other tasks do not access the queue. */
475     vTaskSuspendAll();
476
477     /* Make sure interrupts do not access the queue event list. */
478     prvLockQueue( pxQueue );
479
480     /* If the queue is already full we may have to block. */
481     if( prvIsQueueFull( pxQueue ) )
482     {
483         /* The queue is full - do we want to block or just leave without
484         posting? */
485         if( xTicksToWait > ( portTickType ) 0 )
486         {
487             /* We are going to place ourselves on the xTasksWaitingToSend event list, and will get woken should
488             the delay expire, or space become available on the queue. As detailed above we do not require mutual
489             exclusion on the event list as nothing else can modify it or the ready lists while we have the
490             scheduler suspended and queue locked.
491
492             It is possible that an ISR has removed data from the queue since we checked if any was available. If
493             this is the case then the data will have been copied from the queue, and the queue variables updated,
494             but the event list will not yet have been checked to see if anything is waiting as the queue is
495             locked. */
496             vTaskPlaceOnEventList( &(amp; pxQueue->xTasksWaitingToSend ), xTicksToWait );
497
498             /* Force a context switch now as we are blocked. We can do this from within a critical section as the
499             task we are switching to has its own context. When we return here (i.e. we unblock) we will leave the
500             critical section as normal.
501
502             It is possible that an ISR has caused an event on an unrelated and unlocked queue. If this was the
503             case then the event list for that queue will have been updated but the ready lists left unchanged -
504             instead the readied task will have been added to the pending ready list. */
505             taskENTER_CRITICAL();
506             {
507                 /* We can safely unlock the queue and scheduler here as interrupts are disabled. We must not yield
508                 with anything locked, but we can yield from within a critical section.
509
510                 Tasks that have been placed on the pending ready list cannot be tasks that are waiting for events on
511                 this queue. See in comment xTaskRemoveFromEventList(). */
512                 prvUnlockQueue( pxQueue );
513
514                 /* Resuming the scheduler may cause a yield. If so then there
515                 is no point yielding again here. */
516                 if( !xTaskResumeAll() )
517                 {
518                     taskYIELD();
519                 }
520             }
521
522             /* Before leaving the critical section we have to ensure exclusive access again. */
523             vTaskSuspendAll();
524             prvLockQueue( pxQueue );
525         }
526         taskEXIT_CRITICAL();
527     }
528 }
529
530 /* When we are here it is possible that we unblocked as space became available on the queue.
531 It is also possible that an ISR posted to the queue since we left the critical section, so it may be
532 that again there is no space. This would only happen if a task and ISR post onto the same queue. */
533 taskENTER_CRITICAL();
534 {

```



```

535     if( pxQueue->uxMessagesWaiting < pxQueue->uxLength )
536     {
537         /* There is room in the queue, copy the data into the queue. */
538         prvCopyQueueData( pxQueue, pvItemToQueue );
539         xReturn = pdPASS;
540
541         /* Update the TxLock count so prvUnlockQueue knows to check for
542         tasks waiting for data to become available in the queue. */
543         ++( pxQueue->xTxLock );
544     }
545     else
546     {
547         xReturn = errQUEUE_FULL;
548     }
549 }
550 taskEXIT_CRITICAL();
551
552 /* We no longer require exclusive access to the queue. prvUnlockQueue will remove any tasks suspended
553 on a receive if either this function or an ISR has posted onto the queue. */
554 if( prvUnlockQueue( pxQueue ) )
555 {
556     /* Resume the scheduler - making ready any tasks that were woken by an event while the scheduler was
557     locked. Resuming the scheduler may cause a yield, in which case there is no point yielding again
558     here. */
559     if( !xTaskResumeAll() )
560     {
561         taskYIELD();
562     }
563 }
564 else
565 {
566     /* Resume the scheduler - making ready any tasks that were woken
567     by an event while the scheduler was locked. */
568     xTaskResumeAll();
569 }
570
571 return xReturn;
572 }
573
574 signed portBASE_TYPE xQueueSendFromISR( xQueueHandle pxQueue, const void *pvItemToQueue, signed portBASE_TYPE
575 xTaskPreviouslyWoken )
576 {
577     /* Similar to xQueueSend, except we don't block if there is no room in the queue. Also we don't
578     directly wake a task that was blocked on a queue read, instead we return a flag to say whether a
579     context switch is required or not (i.e. has a task with a higher priority than us been woken by this
580     post). */
581     if( pxQueue->uxMessagesWaiting < pxQueue->uxLength )
582     {
583         prvCopyQueueData( pxQueue, pvItemToQueue );
584
585         /* If the queue is locked we do not alter the event list. This will
586         be done when the queue is unlocked later. */
587         if( pxQueue->xTxLock == queueUNLOCKED )
588         {
589             /* We only want to wake one task per ISR, so check that a task has
590             not already been woken. */
591             if( !xTaskPreviouslyWoken )
592             {
593                 if( !listLIST_IS_EMPTY( &(amp; pxQueue->xTasksWaitingToReceive) ) )
594                 {
595                     if( xTaskRemoveFromEventList( &(amp; pxQueue->xTasksWaitingToReceive) )
596                     != pdFALSE )
597                     {
598                         /* The task waiting has a higher priority so record that a
599                         context switch is required. */
600                         return pdTRUE;
601                     }
602                 }
603             }
604         }
605         else
606         {
607             /* Increment the lock count so the task that unlocks the queue
608             knows that data was posted while it was locked. */
609             ++( pxQueue->xTxLock );
610         }
611     }
612
613     return xTaskPreviouslyWoken;
614 }

```

```

617 signed portBASE_TYPE xQueueReceive( xQueueHandle pxQueue, void *pvBuffer, portTickType xTicksToWait )
618 {
619     signed portBASE_TYPE xReturn;
620
621     /* This function is very similar to xQueueSend(). See comments within
622     xQueueSend() for a more detailed explanation.*/
623
624     /* Make sure other tasks do not access the queue. */
625     vTaskSuspendAll();
626
627     /* Make sure interrupts do not access the queue. */
628     prvLockQueue( pxQueue );
629
630     /* If there are no messages in the queue we may have to block. */
631     if( prvIsQueueEmpty( pxQueue ) )
632     {
633         /* There are no messages in the queue, do we want to block or just leave with nothing? */
634         if( xTicksToWait > ( portTickType ) 0 )
635         {
636             vTaskPlaceOnEventList( &(amp; pxQueue->xTasksWaitingToReceive ), xTicksToWait );
637             taskENTER_CRITICAL();
638             {
639                 prvUnlockQueue( pxQueue );
640                 if( !xTaskResumeAll() )
641                 {
642                     taskYIELD();
643                 }
644
645                 vTaskSuspendAll();
646                 prvLockQueue( pxQueue );
647             }
648             taskEXIT_CRITICAL();
649         }
650     }
651
652     taskENTER_CRITICAL();
653     {
654         if( pxQueue->uxMessagesWaiting > ( unsigned portBASE_TYPE ) 0 )
655         {
656             pxQueue->pcReadFrom += pxQueue->uxItemSize;
657             if( pxQueue->pcReadFrom >= pxQueue->pcTail )
658             {
659                 pxQueue->pcReadFrom = pxQueue->pcHead;
660             }
661             --( pxQueue->uxMessagesWaiting );
662             memcpy( ( void * ) pvBuffer, ( void * ) pxQueue->pcReadFrom,
663                 ( unsigned ) pxQueue->uxItemSize );
664
665             /* Increment the lock count so prvUnlockQueue knows to check for
666             tasks waiting for space to become available on the queue. */
667             ++( pxQueue->xRxLock );
668             xReturn = pdPASS;
669         }
670         else
671         {
672             xReturn = pdFAIL;
673         }
674     }
675     taskEXIT_CRITICAL();
676
677     /* We no longer require exclusive access to the queue. */
678     if( prvUnlockQueue( pxQueue ) )
679     {
680         if( !xTaskResumeAll() )
681         {
682             taskYIELD();
683         }
684     }
685     else
686     {
687         xTaskResumeAll();
688     }
689
690     return xReturn;
691 }
692
693

```

```

694 signed portBASE_TYPE xQueueReceiveFromISR( xQueueHandle pxQueue, void *pvBuffer, signed portBASE_TYPE
695 *pxTaskWoken )
696 {
697     signed portBASE_TYPE xReturn;
698
699     /* We cannot block from an ISR, so check there is data available. */
700     if( pxQueue->uxMessagesWaiting > ( unsigned portBASE_TYPE ) 0 )
701     {
702         /* Copy the data from the queue. */
703         pxQueue->pcReadFrom += pxQueue->uxItemSize;
704         if( pxQueue->pcReadFrom >= pxQueue->pcTail )
705         {
706             pxQueue->pcReadFrom = pxQueue->pcHead;
707         }
708         --( pxQueue->uxMessagesWaiting );
709         memcpy( ( void * ) pvBuffer, ( void * ) pxQueue->pcReadFrom,
710             ( unsigned ) pxQueue->uxItemSize );
711
712         /* If the queue is locked we will not modify the event list. Instead we update the lock count
713         so the task that unlocks the queue will know that an ISR has removed data while the queue was
714         locked. */
715         if( pxQueue->uxRxLock == queueUNLOCKED )
716         {
717             /* We only want to wake one task per ISR, so check that a task has not already been woken. */
718             if( !( *pxTaskWoken ) )
719             {
720                 if( !listLIST_IS_EMPTY( &(amp; pxQueue->xTasksWaitingToSend) ) )
721                 {
722                     if( xTaskRemoveFromEventList( &(amp; pxQueue->xTasksWaitingToSend) )
723                         != pdFALSE )
724                     {
725                         /* The task waiting has a higher priority than us so
726                         force a context switch. */
727                         *pxTaskWoken = pdTRUE;
728                     }
729                 }
730             }
731         }
732         else
733         {
734             /* Increment the lock count so the task that unlocks the queue
735             knows that data was removed while it was locked. */
736             ++( pxQueue->uxRxLock );
737         }
738
739         xReturn = pdPASS;
740     }
741     else
742     {
743         xReturn = pdFAIL;
744     }
745
746     return xReturn;
747 }
748
749 unsigned portBASE_TYPE uxQueueMessagesWaiting( xQueueHandle pxQueue )
750 {
751     unsigned portBASE_TYPE uxReturn;
752
753     taskENTER_CRITICAL();
754     uxReturn = pxQueue->uxMessagesWaiting;
755     taskEXIT_CRITICAL();
756
757     return uxReturn;
758 }
759
760 void vQueueDelete( xQueueHandle pxQueue )
761 {
762     vPortFree( pxQueue->pcHead );
763     vPortFree( pxQueue );
764 }
765
766

```

```

767 static signed portBASE_TYPE prvUnlockQueue( xQueueHandle pxQueue )
768 {
769     signed portBASE_TYPE xYieldRequired = pdFALSE;
770
771     /* THIS FUNCTION MUST BE CALLED WITH THE SCHEDULER SUSPENDED. */
772
773     /* The lock counts contains the number of extra data items placed or
774     removed from the queue while the queue was locked. When a queue is
775     locked items can be added or removed, but the event lists cannot be
776     updated. */
777     taskENTER_CRITICAL();
778     {
779         --( pxQueue->xTxLock );
780
781         /* See if data was added to the queue while it was locked. */
782         if( pxQueue->xTxLock > queueUNLOCKED )
783         {
784             pxQueue->xTxLock = queueUNLOCKED;
785
786             /* Data was posted while the queue was locked. Are any tasks
787             blocked waiting for data to become available? */
788             if( !listLIST_IS_EMPTY( &(amp; pxQueue->xTasksWaitingToReceive) ) )
789             {
790                 /* Tasks that are removed from the event list will get added to
791                 the pending ready list as the scheduler is still suspended. */
792                 if( xTaskRemoveFromEventList( &(amp; pxQueue->xTasksWaitingToReceive) ) != pdFALSE )
793                 {
794                     /* The task waiting has a higher priority so record that a
795                     context switch is required. */
796                     xYieldRequired = pdTRUE;
797                 }
798             }
799         }
800     }
801     taskEXIT_CRITICAL();
802
803     /* Do the same for the Rx lock. */
804     taskENTER_CRITICAL();
805     {
806         --( pxQueue->xRxLock );
807
808         if( pxQueue->xRxLock > queueUNLOCKED )
809         {
810             pxQueue->xRxLock = queueUNLOCKED;
811
812             if( !listLIST_IS_EMPTY( &(amp; pxQueue->xTasksWaitingToSend) ) )
813             {
814                 if( xTaskRemoveFromEventList( &(amp; pxQueue->xTasksWaitingToSend) ) != pdFALSE )
815                 {
816                     xYieldRequired = pdTRUE;
817                 }
818             }
819         }
820     }
821     taskEXIT_CRITICAL();
822
823     return xYieldRequired;
824 }
825
826
827 static signed portBASE_TYPE prvIsQueueEmpty( const xQueueHandle pxQueue )
828 {
829     signed portBASE_TYPE xReturn;
830
831     taskENTER_CRITICAL();
832     xReturn = ( pxQueue->uxMessagesWaiting == ( unsigned portBASE_TYPE ) 0 );
833     taskEXIT_CRITICAL();
834
835     return xReturn;
836 }
837
838 static signed portBASE_TYPE prvIsQueueFull( const xQueueHandle pxQueue )
839 {
840     signed portBASE_TYPE xReturn;
841
842     taskENTER_CRITICAL();
843     xReturn = ( pxQueue->uxMessagesWaiting == pxQueue->uxLength );
844     taskEXIT_CRITICAL();
845
846     return xReturn;
847 }

```

Solutions 2009

FOUR Questions in 180 minutes => 45 min per question

Answer codes: A=analysis, B=bookwork, D=design, C= new application of learnt theory

Question 1.

This question tests whether the students understand some of the issues when writing application code for RTOS

a)

Deadlock:

- cyclic resource dependency graph.
- Once started cannot end regardless of schedule.
- Must involve 2 or more tasks indefinitely suspended

Starvation:

- 100% utilisation of some resource (e.g. CPU)
- can be stopped at any time by changing schedule
- 1 or more tasks indefinitely suspended

Livelock:

- All tasks execute but no progress is made by one or more tasks
- dependent on relative timing between tasks

[8B]

b)

TaskA()

```
{
  done = 0
  do something
  while (not done) {
    lockA:= 1;
    delay(1)
    if (lockB == 0) {
      do something
      done = 1;
      lockA = 0;
    }
  }
}
```

TaskB identical but swap A,B

[8C]

c)

Pseudo-livelock (AKA pseudo-deadlock) where resources are used with a cyclic dependency graph but busy-wait means that no process suspends. The situation is effectively deadlock.

[4A]

Question 2.

a)

1. Y: *QueueReceive()* (API call in Y causes Y to block)
2. X: *QueueSend()* (API call in X causes Y to go from blocked to ready)
3. X: *QueueSend()* (API call in X causes X to block)
4. Y: starts executing (caused by call to scheduler when no higher priority task is ready)
5. X: Y executing causes the message to be picked up and X to unblock and post its second message

[5A]

b)

1. 619-642
2. 473-482, 535-544, 550-559, 561?, 571
3. 472-519
4. 643-669, 675-678, 680 or 680-682 or 687, 690
5. 520-544, 550-554, 571

[5B/A]

c)

QueueReceive() returns an error code if called with no message available and then a task posts a message which unblocks the task, however before it executes another task calls *QueueReceive()* and picks up the posted message.

[5B/A]

d)

Ideally, at the time that no higher priority task is executing, the highest priority waiting task on a queue should receive a posted message. The problem is that when a message is posted a higher priority task may be executing thereby delaying the time at which this decision should be made. But if it is not known which task to make ready until the scheduler is called at some subsequent time this greatly complicates the RTOS – this condition would need to be checked for every time a task suspended.

[5B]

Question 3

This question tests student's understanding of scheduling algorithms and analysis of real-time performance. The last two parts require innovative thinking.

a)

(i) Use RMA priority order, longer periods have lower priorities

(ii)

$$U = 20/200 + 50/500 + 100/t_2 + 50/t_1 = 0.2 + 100/t_2 + 50/t_1 < 4(2^{1/4} - 1) = 0.757 \Rightarrow 100/t_2 + 50/t_1 < 0.557$$

For U meeting this inequality C applies.

For U > 100% A applies

Otherwise B applies.

[8A]

b)

W is higher priority than Z therefore Z can block W but W cannot block Z. Blocking time is therefore 20us max per W period. Using extended RMA we add $20/200=0.1$ onto the utilisation to get new inequality:

$$100/t_2 + 50/t_1 < 0.457$$

Note that X,Y have no extra blocking time. The affect of W being blocked on all other tasks is accounted for in the increased utilisation from W.

[6A]

c)

The maximum time W cannot execute because Z is executing (per W period) is the job time of Z (50us). So the utilisation becomes:

$$(50+20)/200 + 50/500 + 100/t_2 + 50/t_1 = 0.45 + 100/t_2 + 50/t_1 \Rightarrow 100/t_2 + 50/t_1 < 0.307$$

[6A]

Question 4.

This question requires synchronisation similar to barrier – but with the two sides of the barrier separated in time. Many solutions are possible. One solution is:

```
xSemaphoreHandle  SemA1, SemB1, SemC1, SemX2;

initialise()
{
    vSemaphoreBinaryCreate(SemA1) ;
    xSemaphoreTake(SemA1); /* start with no token */
    /* repeat for all other semaphores */
}
SD1()
{
    /* this code will wait till A,B,C have executed printer_open()
    xSemaphoreTake( SemA1);
    xSemaphoreTake(SemaB1);
    xSemaphoreTake(SemaC1);
}
SD2()
{
    int i;
    for (i=0; i < 3; i++) SemaGive(SemaX2);
    /* this code will release A,B,C to execute printer_close()
    xSemaphoreTake( SemaA1);
    xSemaphoreTake(SemaB1);
    xSemaphoreTake(SemaC1);
    for (i=0; i < 3; i++) SemaGive(SemaX2);
    /* this code will release A,B,C to execute printer_close()
}
SX1()
{
    SemaGive(SemaX1); /* must precede print() */
    SemaTake(SemaX2); /* succeeds after print() */
}
SX2()
{
    SemaGive(SemaX1);
    SemaTake(SemaX2);
}
```

(3) print() can only be executed when D has all three tokens and therefore A,B,C are inside SX1 or SX2

(1,2) From initialisation all printer_open() calls must happen before D finishes SD1, whereas all printer_close() calls must happen after print() but before any task finished SX2.

[20D]

Question 5.

a)
`move(int x , int y) {
 lock scheduler
 a := B(x);
 b := B(y);
 B(x) := f(a,b);
 B(y) := f(b,a);
 unlock scheduler
}`

[6D]

b) Use an array of binary semaphores `Sem[i]` , one semaphore for each rod, each initialised to a token count of 1.

```
move( int x, int y) {  
  if (x > y) {  
    SemaTake(Sem[x]);  
    SemaTake(Sem[y]);  
  } else {  
    SemaTake(Sem[y]);  
    SemaTake(Sem[x]);  
  }  
  a := B(x);  
  b := B(y);  
  B(x) := f(a,b);  
  B(y) := f(b,a);  
  SemaGive(Sem[x]);  
  SemaGive(Sem[y]);  
}
```

Semaphore better - independent MOVE operations do not slow each other

Worise - slower + more RAM used.

[10D]

c)

The order in which semaphore `Sem[x]` is claimed is always larger x first hence cyclic dependence graph can never be formed and deadlock not happen.

[4A]

Question 6.

(a)

Interrupt locking temporarily stops interrupts from happening. Any outstanding interrupts are executed at end of section.

Scheduler locking prevents pre-emption from happening but allows interrupts. The highest priority ready task will be scheduled at the end of the section.

[4B]

(b)

Interrupt lock. If CS length is greater than current max CS length it will increase interrupt latency and (therefore) also task latency.

Scheduler lock. If lock length is greater than current max scheduler lock length it will increase task latency, but have no effect on interrupt latency.

Semaphore protection. No effect on task or interrupt latency.

[6A/B]

(c)

(i) If interrupts are locked the tick ISR is delayed until end of lock, and then happens as normal.

(ii) If scheduler is locked the tick ISR happens but does nothing except:

Call user hook function (if defined)

add one to uxTicksMissed (this global is 0 when scheduler is unlocked)

Normally the tick ISR executed the tick function, see below.

When the scheduler is unlocked, if uxTicksMissed is >0 the tick function is executed a number of times equal to the value of uxTicksMissed, and uxTicksMissed set back to 0.

The tick function itself executes:

increment xTickCount

if xTickCount is 0 (therefore has overflowed) swap DelayedTaskList and OverflowDelayedTaskList

Scan through DelayedTaskList and remove tasks that have timeout == xTickCount, adding them to ready list.

Move tasks from PendingTaskList to DelayedTaskList

[10A]