

DEPARTMENT OF ELECTRICAL AND ELECTRONIC ENGINEERING
EXAMINATIONS 2010

EEE/ISE PART III/IV: MEng, BEng and ACGI

REAL-TIME OPERATING SYSTEMS

Thursday, 13 May 10:00 am

Time allowed: 1:30 hours

There are **SIX** questions on this paper.

Answer **FOUR** questions.

All questions carry equal marks.

Any special instructions for invigilators and information for candidates are on page 1.

Examiners responsible First Marker(s) : T.J.W. Clarke
 Second Marker(s) : Y.K. Demiris

Special Information for Invigilators: none.

Information for Candidates

Pre-published course notes can be found in the booklet RTOS Exam Notes.

The Questions

1.

- a) State the assumptions behind the job model of computation, and the Rate Monotonic Analysis (RMA) utilisation limit. [4]
- b) Table 1.1 shows the CPU time (in cycles) and period (in microseconds) of four jobs. Using rate monotonic analysis or otherwise determine the optimum priorities and calculate a limit on the CPU cycle time which will ensure all deadlines are met if jobs have optimal priorities [5]
- c) Assume all four jobs from part b) use a single mutually exclusive shared resource once per job time. If the time for which each job locks the resource is R CPU cycles, and priority inversion is prevented, by usages of priority inheritance protocol, determine as a function of R the blocking time of each job for use in extended RMA, and hence calculate a new limit on CPU cycle time for all deadlines to be met. [5]
- d) The jobs from part b) are implemented so that the two lowest priority jobs have equal priority, but priorities are otherwise as in part b). Assume R is zero (no time waiting on shared resource). If no assumptions are made about which job is executed first when both jobs have equal priority, compare the execution of this system with the same system under optimal priorities. Hence, or otherwise, determine the correct blocking time for each job in order to apply extended RMA to the three priority system, and therefore calculate an inequality on the CPU cycle time for all deadlines to be met. [6]

Job	CPU time / cycles	Period / μ s
A	100	10
B	200	40
C	300	150
D	400	100

Table 1.1

- 2.
- a) Describe the main performance characteristics of the FreeRTOS double-linked list package, and list the places in which it is used in FreeRTOS. [8]
 - b) Explain the use of sentinel nodes in the FreeRTOS doubly-linked list package. [4]
 - c) For each usage in part a), compare and contrast the use of doubly-linked lists with a similar singly-linked list package, using circular lists and a sentinel node. [8]
- 3.
- a) Explain how queues can be used to implement semaphores. [4]
 - b) The pseudocode in Figure 3.1 describes a real-time system with tasks A, B, C, D and corresponding pseudocode **TaskA**, **TaskB**, **TaskC**, **TaskD**. The two queues **q[1]** and **q[2]** are initialised by the code **Init()**. The pseudocode comments *execute* and *wait* represent periods when the task is respectively executing, or suspended, for fixed but unspecified lengths of time.
Identify all potential liveness problems, stating the conditions required for each problem. [16]

```

Init()
{
  for (i = 1; i < 3; i++) {
    q[i] = QueueCreate();
    QueueSend(q[i]);
  }
}

TaskA()
{
  while (TRUE) {
    QueueReceive(q[2]);
    QueueReceive(q[1]);
    /* execute */
    QueueSend(q[2]);
    QueueSend(q[1]);
    /* wait */
  }
}

TaskB()
{
  while (TRUE) {
    QueueReceive(q[1]);
    QueueReceive(q[2]);
    /* execute */
    QueueSend(q[2]);
    QueueSend(q[1]);
    /* wait */
  }
}

TaskC()
{
  while(TRUE) {
    QueueReceive(q[1]);
    /* execute */
    QueueSend(q[1]);
    /* wait */
  }
}

TaskD()
{
  while(TRUE) {
    /* execute */
    /* wait */
  }
}

```

Figure 3.1

4.

- a) Describe the three types of locking used in the FreeRTOS queue package. [5]
- b) Determine, with reference to the line numbers shown in the Exam Notes, at which points in the execution of QueueReceive() locking is changed. In each case, where a lock is started or stopped, state why this is either necessary or desirable. [15]

5.

Table 5.1 Lists the I/O devices serviced by a microcontroller, together with the required maximum latency time, frequency at which the device needs service, and length of service routine. Assume that interrupt entry time is $E \mu\text{s}$ and return time is $R \mu\text{s}$, and that two interrupt levels are available $P1$ and $P2$ with $P1$ higher priority than $P2$ so that each device may raise a $P1$ or $P2$ interrupt, or no interrupt. Each interrupt has a separate vector, however when more than one device can raise an interrupt both use the same vector.

- a) Explain how a foreground/background system can be used to implement the system, and why the latency time of any device can be delayed by at most one service of any other device. [5]
- b) Suppose two devices A & B are serviced from the same interrupt. Determine from the polling code in Table 5.2 what is the maximum latency for each device as a function of E , R , and device service times, assuming no other interrupts or critical sections. [5]
- c) Assume $E = R = 0$, and that polling overhead is negligible. List all possible strategies for servicing the devices which meet all latency times. [5]
- d) For each strategy in part c), determine the inequalities on E , R which must be satisfied for the strategy to meet each latency time, again assuming that polling time is negligible. [5]

Device	Service Time $S / \mu\text{s}$	Service Frequency / Hz	Maximum Latency Time $L / \mu\text{s}$
keyboard	10	100	90
printer	10	2000	15
network	1	10000	3
disk	1	4000	1.5

Table 5.1

```
if (device A ready) service device A;  
if (device B ready) service device B;
```

Table 5.2

6. A digital oscilloscope (DSO) front-end receives real-time data samples from four independent channels at possibly different data rates, and processes these samples independently with a fast fourier transform (FFT) to generate phase and magnitude spectra. The spectra are further processed to generate time-averaged spectra, which are sent in real time to four separate windows on an LCD screen.

The DSO can simultaneously display and output averaged spectra via a serial link which sends 8 bit bytes at a rate of 1000 bytes per second.

Operation modes are controlled by a keyboard which must be polled once every 10ms in order to detect key presses. Keys may affect all other operations.

A single hardware FFT unit is available which calculates one FFT in the time shown in Table 6.1 and is used to implement all the FFT operations. The CPU time used is the sum of that needed to transfer input and output to or from the unit, and is in addition to the FFT calculation time during which the CPU is not required.

Each channel may use a 256, 1024, or 4096 sample FFT as selected by the user, however at higher sample rates and large FFT sizes there will be a gap between successive samples due to the time taken by the FFT.

- a) Using the outside-in design strategy or otherwise, determine an appropriate set of tasks to implement this system, stating the function of each task.

[6]

- b) Considering CPU and FFT unit utilisation, and stating any assumptions you make, calculate the fastest sample rate which can be processed with 4096 sample FFTs on all channels without missing any samples.

[8]

- c) Discuss how the FFT unit should interact with a real-time operating system, how it is protected from incorrect use, and whether the system is likely to suffer liveness problems.

[6]

Number of samples	FFT time	CPU time
256	10 μ s	5 μ s
1024	60 μ s	20 μ s
4096	260 μ s	80 μ s

Table 6.1

RTOS EXAM NOTES 2010

Task.h

```

1 1
2 2
3 3  typedef void * xTaskHandle;
4 4  #define taskYIELD()          portYIELD()
5 5  #define taskENTER_CRITICAL()  portENTER_CRITICAL()
6 6  #define taskEXIT_CRITICAL()   portEXIT_CRITICAL()
7 7  #define taskDISABLE_INTERRUPTS() portDISABLE_INTERRUPTS()
8 8  #define taskENABLE_INTERRUPTS() portENABLE_INTERRUPTS()
9 9

10 /*-----
11  *  TASK CREATION API
12  *-----*/
13
14 signed portBASE_TYPE xTaskCreate(      pdTASK_CODE pvTaskCode, const signed portCHAR * const pcName,
15                                     unsigned portSHORT usStackDepth, void *pvParameters,
16                                     unsigned portBASE_TYPE uxPriority, xTaskHandle *pvCreatedTask );
17
18 void vTaskDelete( xTaskHandle pxTask );
19
20 /*-----
21  *  TASK CONTROL API
22  *-----*/
23
24 void vTaskDelay( portTickType xTicksToDelay );
25 void vTaskDelayUntil( portTickType *pxPreviousWakeTime, portTickType xTimeIncrement );
26 unsigned portBASE_TYPE uxTaskPriorityGet( xTaskHandle pxTask );
27 void vTaskPrioritySet( xTaskHandle pxTask, unsigned portBASE_TYPE uxNewPriority );
28 void vTaskSuspend( xTaskHandle pxTaskToSuspend );
29 void vTaskResume( xTaskHandle pxTaskToResume );
30 portBASE_TYPE xTaskResumeFromISR( xTaskHandle pxTaskToResume );
31
32 /*-----
33  *  SCHEDULER CONTROL
34  *-----*/
35
36 void vTaskStartScheduler( void );
37 void vTaskEndScheduler( void );
38 void vTaskSuspendAll( void );
39 signed portBASE_TYPE xTaskResumeAll( void );
40
41 /*-----
42  *  TASK UTILITIES
43  *-----*/
44
45 portTickType xTaskGetTickCount( void );
46 unsigned portBASE_TYPE uxTaskGetNumberOfTasks( void );
47 void vTaskPlaceOnEventList( xList *pxEventList, portTickType xTicksToWait );
48 signed portBASE_TYPE xTaskRemoveFromEventList( const xList *pxEventList );
49 void vTaskCleanUpResources( void );
50 inline void vTaskSwitchContext( void );
51 xTaskHandle xTaskGetCurrentTaskHandle( void );
52
53
54
55 Semaphr.c
56
57 #define vSemaphoreCreateBinary( xSemaphore )           \
58     xSemaphore = xQueueCreate( ( unsigned portCHAR ) 1, semSEMAPHORE_QUEUE_ITEM_LENGTH ); \
59     if( xSemaphore != NULL )                      \
60     {                                              \
61         xSemaphoreGive( xSemaphore );              \
62     }                                              \
63 }
64
65 #define xSemaphoreTake( xSemaphore, xBlockTime )        \
66     xQueueReceive( ( xQueueHandle ) xSemaphore, NULL, xBlockTime )
67
68 #define xSemaphoreGive( xSemaphore ) xQueueSend( ( xQueueHandle ) xSemaphore, NULL, semGIVE_BLOCK_TIME )
69
70 #define xSemaphoreGiveFromISR( xSemaphore, xTaskPreviouslyWoken ) \
71     xQueueSendFromISR( ( xQueueHandle ) xSemaphore, NULL, xTaskPreviouslyWoken )

```

```

100 From Task.h - related to lists package
101 signed portBASE_TYPE xTaskRemoveFromEventList( const xList *pxEventList )
102 {
103     tskTCB *pxUnblockedTCB;
104     portBASE_TYPE xReturn;
105
106     /* THIS FUNCTION MUST BE CALLED WITH INTERRUPTS DISABLED OR THE
107      SCHEDULER SUSPENDED. It can also be called from within an ISR. */
108
109     /* The event list is sorted in priority order, so we can remove the
110      first in the list, remove the TCB from the delayed list, and add
111      it to the ready list.
112
113     If an event is for a queue that is locked then this function will never
114     get called - the lock count on the queue will get modified instead. This
115     means we can always expect exclusive access to the event list here. */
116     pxUnblockedTCB = ( tskTCB * ) listGET_OWNER_OF_HEAD_ENTRY( pxEventList );
117     vListRemove( &( pxUnblockedTCB->xEventListItem ) );
118
119     if( uxSchedulerSuspended == ( unsigned portBASE_TYPE ) pdFALSE )
120     {
121         vListRemove( &( pxUnblockedTCB->xGenericListItem ) );
122         prvAddTaskToReadyQueue( pxUnblockedTCB );
123     }
124     else
125     {
126         /* We cannot access the delayed or ready lists, so will hold this
127          task pending until the scheduler is resumed. */
128         vListInsertEnd( ( xList * ) &( xPendingReadyList ), &( pxUnblockedTCB->xEventListItem ) );
129     }
130
131     if( pxUnblockedTCB->uxPriority >= pxCurrentTCB->uxPriority )
132     {
133         /* Return true if the task removed from the event list has
134          a higher priority than the calling task. This allows
135          the calling task to know if it should force a context
136          switch now. */
137         xReturn = pdTRUE;
138     }
139     else
140     {
141         xReturn = pdFALSE;
142     }
143
144     return xReturn;
145 }

```

```

146 List.h
147 /*
148  * Definition of the only type of object that a list can contain.
149  */
150 struct xLIST_ITEM
151 {
152     portTickType xItemValue;           /*< The value being listed. In most cases this is
153                                         used to sort the list in descending order. */
154     volatile struct xLIST_ITEM * pxNext; /*< Pointer to the next xListItem in the list. */
155     volatile struct xLIST_ITEM * pxPrevious; /*< Pointer to the previous xListItem in the list. */
156     void * pvOwner;                  /*< Pointer to the object (normally a TCB) that contains the list item. */
157     void * pvContainer;             /*< Pointer to the list in which this list item is placed (if any). */
158 };
159 typedef struct xLIST_ITEM xListItem; /* For some reason lint wants this as two separate definitions. */
160
161 struct xMINI_LIST_ITEM
162 {
163     portTickType xItemValue;
164     volatile struct xLIST_ITEM *pxNext;
165     volatile struct xLIST_ITEM *pxPrevious;
166 };
167 typedef struct xMINI_LIST_ITEM xMiniListItem;
168
169 /*
170  * Definition of the type of queue used by the scheduler.
171  */
172 typedef struct xLIST
173 {
174     volatile unsigned portBASE_TYPE uxNumberOfItems;
175     volatile xListItem * pxIndex;        /* Used to walk through the list */
176     volatile xMinilistItem xListEnd;    /* List item that contains the maximum possible item value */
177 } xList;
178
179 #define listSET_LIST_ITEM_OWNER( pxListItem, pxOwner ) ( pxListItem )->pvOwner = ( void * ) pxOwner
180
181 #define listSET_LIST_ITEM_VALUE( pxListItem, xValue )          ( pxListItem )->xItemValue = xValue
182
183 #define listGET_LIST_ITEM_VALUE( pxListItem )                 ( ( pxListItem )->xItemValue )
184
185 #define listLIST_IS_EMPTY( pxList )   ( ( pxList )->uxNumberOfItems == ( unsigned portBASE_TYPE ) 0 )
186
187 #define listCURRENT_LIST_LENGTH( pxList )      ( ( pxList )->uxNumberOfItems )
188
189 #define listGET_OWNER_OF_NEXT_ENTRY( pxtcb, pxList ) \
190     /* Increment the index to the next item and return the item, ensuring */ \
191     /* we don't return the marker used at the end of the list. */ \
192     ( pxList )->pxIndex = ( pxList )->pxIndex->pxNext; \
193     if( ( pxList )->pxIndex == ( xListItem * ) &( ( pxList )->xListEnd ) ) \
194     { \
195         ( pxList )->pxIndex = ( pxList )->pxIndex->pxNext; \
196     } \
197     pxtcb = ( pxList )->pxIndex->pvOwner \
198
199
200 #define listGET_OWNER_OF_HEAD_ENTRY( pxList ) ( ( pxList->uxNumberOfItems != ( unsigned portBASE_TYPE ) 0 \
201 ) ? ( ( pxList->xListEnd )->pxNext->pvOwner ) : ( NULL ) )
202
203 #define listIS_CONTAINED_WITHIN( pxList, pxListItem ) ( ( pxListItem )->pvContainer == ( void * ) pxList )
204
205 void vListInitialise( xList *pxList );
206
207 void vListInitialiseItem( xListItem *pxItem );
208
209 void vListInsert( xList *pxList, xListItem *pxNewListItem );
210
211 void vListInsertEnd( xList *pxList, xListItem *pxNewListItem );
212
213 void vListRemove( xListItem *pxItemToRemove );

```

```

214 List.c
215 #include <stdlib.h>
216 #include "FreeRTOS.h"
217 #include "list.h"
218
219 -----
220 * PUBLIC LIST API documented in list.h
221 -----*/
222
223 void vListInitialise( xList *pxList )
224 {
225     /* The list structure contains a list item which is used to mark the end of the list. To initialise
226         the list the list end is inserted as the only list entry. */
227     pxList->pxIndex = ( xListItem * ) &( pxList->xListEnd );
228
229     /* The list end value is the highest possible value in the list to ensure it
230         remains at the end of the list. */
231     pxList->xListEnd.xItemValue = portMAX_DELAY;
232
233     /* The list end next and previous pointers point to itself so we know when the list is empty. */
234     pxList->xListEnd.pxNext = ( xListItem * ) &( pxList->xListEnd );
235     pxList->xListEnd.pxPrevious = ( xListItem * ) &( pxList->xListEnd );
236
237     pxList->uxNumberOfItems = 0;
238 }
239
240 void vListInitialiseItem( xListItem *pxItem )
241 {
242     /* Make sure the list item is not recorded as being on a list. */
243     pxItem->pvContainer = NULL;
244 }
245
246 void vListInsertEnd( xList *pxList, xListItem *pxNewListItem )
247 {
248     volatile xListItem * pxIndex;
249
250     /* Insert a new list item into pxList, but rather than sort the list, makes the new list item the
251 last
252     item to be removed by a call to pvListGetOwnerOfNextEntry. This means it has to be the item
253     pointed to by the pxIndex member. */
254     pxIndex = pxList->pxIndex;
255
256     pxNewListItem->pxNext = pxIndex->pxNext;
257     pxNewListItem->pxPrevious = pxList->pxIndex;
258     pxIndex->pxNext->pxPrevious = ( volatile xListItem * ) pxNewListItem;
259     pxIndex->pxNext = ( volatile xListItem * ) pxNewListItem;
260     pxList->pxIndex = ( volatile xListItem * ) pxNewListItem;
261
262     /* Remember which list the item is in. */
263     pxNewListItem->pvContainer = ( void * ) pxList;
264
265     ( pxList->uxNumberOfItems )++;
266 }
267

```

```

268 void vListInsert( xList *pxList, xListItem *pxNewListItem )
269 {
270     volatile xListItem *pxIterator;
271     portTickType xValueOfInsertion;
272
273     /* Insert the new list item into the list, sorted in ulListItem order. */
274     xValueOfInsertion = pxNewListItem->xItemValue;
275
276     /* If the list already contains a list item with the same item value then the new list item should be
277      placed after it. This ensures that TCB's which are stored in ready lists (all of which have the same
278      ulListItem value) get an equal share of the CPU. However, if the xItemValue is the same as the back
279      marker the iteration loop below will not end. This means we need to guard against this by checking
280      the value first and modifying the algorithm slightly if necessary. */
281     if( xValueOfInsertion == portMAX_DELAY )
282     {
283         pxIterator = pxList->xListEnd.pxPrevious;
284     }
285     else
286     {
287         for( pxIterator = ( xListItem * ) &( pxList->xListEnd );
288              pxIterator->pxNext->xItemValue <= xValueOfInsertion;
289              pxIterator = pxIterator->pxNext )
290         {
291             /* There is nothing to do here, we are just iterating to the wanted insertion position. */
292         }
293     }
294
295     pxNewListItem->pxNext = pxIterator->pxNext;
296     pxNewListItem->pxNext->pxPrevious = ( volatile xListItem * ) pxNewListItem;
297     pxNewListItem->pxPrevious = pxIterator;
298     pxIterator->pxNext = ( volatile xListItem * ) pxNewListItem;
299
300     /* Remember which list the item is in. This allows fast removal of the item later. */
301     pxNewListItem->pvContainer = ( void * ) pxList;
302
303     ( pxList->uxNumberOfItems )++;
304 }
305
306 void vListRemove( xListItem *pxItemToRemove )
307 {
308     xList * pxList;
309     pxItemToRemove->pxNext->pxPrevious = pxItemToRemove->pxPrevious;
310     pxItemToRemove->pxPrevious->pxNext = pxItemToRemove->pxNext;
311
312     /* The list item knows which list it is in. Obtain the list from the list item. */
313     pxList = ( xList * ) pxItemToRemove->pvContainer;
314
315     /* Make sure the index is left pointing to a valid item. */
316     if( pxList->pxIndex == pxItemToRemove )
317     {
318         pxList->pxIndex = pxItemToRemove->pxPrevious;
319     }
320
321     pxItemToRemove->pvContainer = NULL;
322     ( pxList->uxNumberOfItems )--;
323 }
324 /*-----*/

```

```

300 Queue.h
301 typedef void * xQueueHandle;
302
303 xQueueHandle xQueueCreate( unsigned portBASE_TYPE uxQueueLength, unsigned portBASE_TYPE uxItemSize );
304
305 signed portBASE_TYPE xQueueSend( xQueueHandle xQueue, const void * pvItemToQueue, portTickType xTicksToWait );
306
307 signed portBASE_TYPE xQueueReceive( xQueueHandle xQueue, void *pvBuffer, portTickType xTicksToWait );
308
309 unsigned portBASE_TYPE uxQueueMessagesWaiting( xQueueHandle xQueue );
310
311 void vQueueDelete( xQueueHandle xQueue );
312
313 signed portBASE_TYPE xQueueSendFromISR( xQueueHandle pxQueue, const void *pvItemToQueue, signed portBASE_TYPE
314 *xTaskPreviouslyWoken );
315
316 signed portBASE_TYPE xQueueReceiveFromISR( xQueueHandle pxQueue, void *pvBuffer, signed portBASE_TYPE
317 *pxTaskWoken );
318

319 Queue.c
320
321 -----*
322 * PUBLIC LIST API documented in list.h
323 -----*/
324
325 /* Constants used with the cRxLock and cTxLock structure members. */
326 #define queueUNLOCKED ( ( signed portBASE_TYPE ) -1 )
327
328 /*
329  * Definition of the queue used by the scheduler.
330  * Items are queued by copy, not reference.
331  */
332 typedef struct QueueDefinition
333 {
334     signed portCHAR *pcHead; /*< Points to the beginning of the queue storage area. */
335     signed portCHAR *pcTail; /*< Points to the byte at the end of the queue storage area.
336             Once more byte is allocated than necessary to store the queue items,
337             this is used as a marker. */
338
339     signed portCHAR *pcWriteTo; /*< Points to the free next place in the storage area. */
340     signed portCHAR *pcReadFrom; /*< Points to the last place that a queued item was read from. */
341
342     xList xTasksWaitingToSend; /*< List of tasks that are blocked waiting to post onto this queue.
343             Stored in priority order. */
344     xList xTasksWaitingToReceive; /*< List of tasks that are blocked waiting to
345             read from this queue. Stored in priority order. */
346
347     unsigned portBASE_TYPE uxMessagesWaiting; /*< The number of items currently in the queue. */
348     unsigned portBASE_TYPE uxLength; /*< The length of the queue defined as the number
349             of items it will hold, not the number of bytes. */
350     unsigned portBASE_TYPE uxItemSize; /*< The size of each items that the queue will hold. */
351
352     signed portBASE_TYPE xRxLock; /*< Stores the number of items received from the queue
353             (removed from the queue) while the queue was locked.
354             Set to queueUNLOCKED when the queue is not locked. */
355     signed portBASE_TYPE xTxLock; /*< Stores the number of items transmitted to the queue
356             (added to the queue) while the queue was locked.
357             Set to queueUNLOCKED when the queue is not locked. */
358 } xQUEUE;
359 -----*/
360
361 /*
362  * Inside this file xQueueHandle is a pointer to a xQUEUE structure.
363  * To keep the definition private the API header file defines it as a
364  * pointer to void.
365  */
366 typedef xQUEUE * xQueueHandle;
367

```

```

368 /*
369  * Unlocks a queue locked by a call to prvLockQueue. Locking a queue does not
370  * prevent an ISR from adding or removing items to the queue, but does prevent
371  * an ISR from removing tasks from the queue event lists. If an ISR finds a
372  * queue is locked it will instead increment the appropriate queue lock count
373  * to indicate that a task may require unblocking. When the queue is unlocked
374  * these lock counts are inspected, and the appropriate action taken.
375  */
376 static signed portBASE_TYPE prvUnlockQueue( xQueueHandle pxQueue );
377
378 /*
379  * Uses a critical section to determine if there is any data in a queue.
380  *
381  * @return pdTRUE if the queue contains no items, otherwise pdFALSE.
382  */
383 static signed portBASE_TYPE prvIsQueueEmpty( const xQueueHandle pxQueue );
384
385 /*
386  * Uses a critical section to determine if there is any space in a queue.
387  *
388  * @return pdTRUE if there is no space, otherwise pdFALSE;
389  */
390 static signed portBASE_TYPE prvIsQueueFull( const xQueueHandle pxQueue );
391
392 /*
393  * Macro that copies an item into the queue. This is done by copying the item
394  * byte for byte, not by reference. Updates the queue state to ensure it's
395  * integrity after the copy.
396  */
397 #define prvCopyQueueData( pxQueue, pvItemToQueue ) \
398 { \
399     memcpy( ( void * ) pxQueue->pcWriteTo, pvItemToQueue, ( unsigned ) pxQueue->uxItemSize ); \
400     ++( pxQueue->uxMessagesWaiting ); \
401     pxQueue->pcWriteTo += pxQueue->uxItemSize; \
402     if( pxQueue->pcWriteTo >= pxQueue->pcTail ) \
403     { \
404         pxQueue->pcWriteTo = pxQueue->pcHead; \
405     } \
406 }
407
408 /*
409  * Macro to mark a queue as locked. Locking a queue prevents an ISR from accessing the queue event lists.
410  */
411 #define prvLockQueue( pxQueue ) \
412 { \
413     taskENTER_CRITICAL(); \
414     ++( pxQueue->xRxLock ); \
415     ++( pxQueue->xTxLock ); \
416     taskEXIT_CRITICAL(); \
417 }
418
419 *-----*
420 * PUBLIC QUEUE MANAGEMENT API documented in queue.h
421 *-----*/
422 xQueueHandle xQueueCreate( unsigned portBASE_TYPE uxQueueLength, unsigned portBASE_TYPE uxItemSize )
423 {
424     xQUEUE *pxNewQueue;
425     size_t xQueueSizeInBytes;
426
427     /* Allocate the new queue structure. */
428     if( uxQueueLength > ( unsigned portBASE_TYPE ) 0 )
429     {
430         pxNewQueue = ( xQUEUE * ) pvPortMalloc( sizeof( xQUEUE ) );
431         if( pxNewQueue != NULL )
432         {
433             /* Create the list of pointers to queue items. The queue is one byte
434             longer than asked for to make wrap checking easier/faster. */
435             xQueueSizeInBytes = ( size_t ) ( uxQueueLength * uxItemSize ) + ( size_t ) 1;
436
437             pxNewQueue->pcHead = ( signed portCHAR * ) pvPortMalloc( xQueueSizeInBytes );
438             if( pxNewQueue->pcHead != NULL )
439             {
440                 /* Initialise the queue members as described above where the
441                 queue type is defined. */
442                 pxNewQueue->pcTail = pxNewQueue->pcHead + ( uxQueueLength * uxItemSize );
443                 pxNewQueue->uxMessagesWaiting = 0;
444                 pxNewQueue->pcWriteTo = pxNewQueue->pcHead;
445                 pxNewQueue->pcReadFrom = pxNewQueue->pcHead + ( ( uxQueueLength - 1 ) *
446                                         uxItemSize );
447                 pxNewQueue->uxLength = uxQueueLength;
448                 pxNewQueue->uxItemSize = uxItemSize;
449                 pxNewQueue->xRxLock = queueUNLOCKED;
450                 pxNewQueue->xTxLock = queueUNLOCKED;

```

```

451
452     /* Likewise ensure the event queues start with the correct state. */
453     vListInitialise( &( pxNewQueue->xTasksWaitingToSend ) );
454     vListInitialise( &( pxNewQueue->xTasksWaitingToReceive ) );
455
456     return pxNewQueue;
457 }
458 else
459 {
460     vPortFree( pxNewQueue );
461 }
462 }
463 }
464
465 /* Will only reach here if we could not allocate enough memory or no memory
466 was required. */
467 return NULL;
468 }
469
470 signed portBASE_TYPE xQueueSend( xQueueHandle pxQueue, const void *pvItemToQueue, portTickType xTicksToWait )
471 {
472     signed portBASE_TYPE xReturn;
473
474     /* Make sure other tasks do not access the queue. */
475     vTaskSuspendAll();
476
477     /* Make sure interrupts do not access the queue event list. */
478     prvLockQueue( pxQueue );
479
480     /* If the queue is already full we may have to block. */
481     if( prvIsQueueFull( pxQueue ) )
482     {
483         /* The queue is full - do we want to block or just leave without
484         posting? */
485         if( xTicksToWait > ( portTickType ) 0 )
486         {
487             /* We are going to place ourselves on the xTasksWaitingToSend event list, and will get woken should
488             the delay expire, or space become available on the queue. As detailed above we do not require mutual
489             exclusion on the event list as nothing else can modify it or the ready lists while we have the
490             scheduler suspended and queue locked.
491
492             It is possible that an ISR has removed data from the queue since we checked if any was available. If
493             this is the case then the data will have been copied from the queue, and the queue variables updated,
494             but the event list will not yet have been checked to see if anything is waiting as the queue is
495             locked. */
496             vTaskPlaceOnEventList( &( pxQueue->xTasksWaitingToSend ), xTicksToWait );
497
498             /* Force a context switch now as we are blocked. We can do this from within a critical section as the
499             task we are switching to has its own context. When we return here (i.e. we unblock) we will leave the
500             critical section as normal.
501
502             It is possible that an ISR has caused an event on an unrelated and unlocked queue. If this was the
503             case then the event list for that queue will have been updated but the ready lists left unchanged -
504             instead the readied task will have been added to the pending ready list. */
505             taskENTER_CRITICAL();
506         }
507
508         /* We can safely unlock the queue and scheduler here as interrupts are disabled. We must not yield
509         with anything locked, but we can yield from within a critical section.
510
511         Tasks that have been placed on the pending ready list cannot be tasks that are waiting for events on
512         this queue. See in comment xTaskRemoveFromEventList(). */
513         prvUnlockQueue( pxQueue );
514
515         /* Resuming the scheduler may cause a yield. If so then there
516         is no point yielding again here. */
517         if( !xTaskResumeAll() )
518         {
519             taskYIELD();
520         }
521
522         /* Before leaving the critical section we have to ensure exclusive access again. */
523         vTaskSuspendAll();
524         prvLockQueue( pxQueue );
525     }
526     taskEXIT_CRITICAL();
527 }
528
529 /* When we are here it is possible that we unblocked as space became available on the queue.
530 It is also possible that an ISR posted to the queue since we left the critical section, so it may be
531 that again there is no space. This would only happen if a task and ISR post onto the same queue. */
532 taskENTER_CRITICAL();
533

```

```

534     if( pxQueue->uxMessagesWaiting < pxQueue->uxLength )
535     {
536         /* There is room in the queue, copy the data into the queue. */
537         prvCopyQueueData( pxQueue, pvItemToQueue );
538         xReturn = pdPASS;
539
540         /* Update the TxLock count so prvUnlockQueue knows to check for
541          tasks waiting for data to become available in the queue. */
542         ++( pxQueue->xTxLock );
543     }
544     else
545     {
546         xReturn = errQUEUE_FULL;
547     }
548 }
549 taskEXIT_CRITICAL();
550
551 /* We no longer require exclusive access to the queue.  prvUnlockQueue will remove any tasks suspended
552 on a receive if either this function or an ISR has posted onto the queue. */
553 if( prvUnlockQueue( pxQueue ) )
554 {
555     /* Resume the scheduler - making ready any tasks that were woken by an event while the scheduler was
556     locked.  Resuming the scheduler may cause a yield, in which case there is no point yielding again
557     here. */
558     if( !xTaskResumeAll() )
559     {
560         taskYIELD();
561     }
562 }
563 else
564 {
565     /* Resume the scheduler - making ready any tasks that were woken
566     by an event while the scheduler was locked. */
567     xTaskResumeAll();
568 }
569
570 return xReturn;
571 }
572
573 signed portBASE_TYPE xQueueSendFromISR( xQueueHandle pxQueue, const void *pvItemToQueue, signed portBASE_TYPE
574 xTaskPreviouslyWoken )
575 {
576     /* Similar to xQueueSend, except we don't block if there is no room in the queue.  Also we don't
577     directly wake a task that was blocked on a queue read, instead we return a flag to say whether a
578     context switch is required or not (i.e. has a task with a higher priority than us been woken by this
579     post). */
580     if( pxQueue->uxMessagesWaiting < pxQueue->uxLength )
581     {
582         prvCopyQueueData( pxQueue, pvItemToQueue );
583
584         /* If the queue is locked we do not alter the event list.  This will
585         be done when the queue is unlocked later. */
586         if( pxQueue->xTxLock == queueUNLOCKED )
587         {
588             /* We only want to wake one task per ISR, so check that a task has
589             not already been woken. */
590             if( !xTaskPreviouslyWoken )
591             {
592                 if( !listLIST_IS_EMPTY( &( pxQueue->xTasksWaitingToReceive ) ) )
593                 {
594                     if( xTaskRemoveFromEventList( &( pxQueue->xTasksWaitingToReceive ) )
595                         != pdFALSE )
596                     {
597                         /* The task waiting has a higher priority so record that a
598                         context switch is required. */
599                         return pdTRUE;
600                     }
601                 }
602             }
603         }
604     }
605     /* Increment the lock count so the task that unlocks the queue
606     knows that data was posted while it was locked. */
607     ++( pxQueue->xTxLock );
608 }
609
610 }
611
612 return xTaskPreviouslyWoken;
613 }
614

```

```

615 signed portBASE_TYPE xQueueReceive( xQueueHandle pxQueue, void *pvBuffer, portTickType xTicksToWait )
616 {
617     signed portBASE_TYPE xReturn;
618
619     /* This function is very similar to xQueueSend(). See comments within
620      xQueueSend() for a more detailed explanation.*/
621
622     /* Make sure other tasks do not access the queue. */
623     vTaskSuspendAll();
624
625     /* Make sure interrupts do not access the queue. */
626     prvLockQueue( pxQueue );
627
628     /* If there are no messages in the queue we may have to block. */
629     if( prvIsQueueEmpty( pxQueue ) )
630     {
631         /* There are no messages in the queue, do we want to block or just leave with nothing? */
632         if( xTicksToWait > ( portTickType ) 0 )
633         {
634             vTaskPlaceOnEventList( &( pxQueue->xTasksWaitingToReceive ), xTicksToWait );
635             taskENTER_CRITICAL();
636             {
637                 prvUnlockQueue( pxQueue );
638                 if( !xTaskResumeAll() )
639                 {
640                     taskYIELD();
641                 }
642
643                 vTaskSuspendAll();
644                 prvLockQueue( pxQueue );
645             }
646             taskEXIT_CRITICAL();
647         }
648     }
649
650     taskENTER_CRITICAL();
651     {
652         if( pxQueue->uxMessagesWaiting > ( unsigned portBASE_TYPE ) 0 )
653         {
654             pxQueue->pcReadFrom += pxQueue->uxItemSize;
655             if( pxQueue->pcReadFrom >= pxQueue->pcTail )
656             {
657                 pxQueue->pcReadFrom = pxQueue->pcHead;
658             }
659             --( pxQueue->uxMessagesWaiting );
660             memcpy( ( void * ) pvBuffer, ( void * ) pxQueue->pcReadFrom,
661                     ( unsigned ) pxQueue->uxItemSize );
662
663             /* Increment the lock count so prvUnlockQueue knows to check for
664              tasks waiting for space to become available on the queue. */
665             ++( pxQueue->xRxLock );
666             xReturn = pdPASS;
667         }
668         else
669         {
670             xReturn = pdFAIL;
671         }
672     }
673     taskEXIT_CRITICAL();
674
675     /* We no longer require exclusive access to the queue. */
676     if( prvUnlockQueue( pxQueue ) )
677     {
678         if( !xTaskResumeAll() )
679         {
680             taskYIELD();
681         }
682     }
683     else
684     {
685         xTaskResumeAll();
686     }
687
688     return xReturn;
689 }
690

```

```

691 signed portBASE_TYPE xQueueReceiveFromISR( xQueueHandle pxQueue, void *pvBuffer, signed portBASE_TYPE
692 *pxTaskWoken )
693 {
694     signed portBASE_TYPE xReturn;
695
696     /* We cannot block from an ISR, so check there is data available. */
697     if( pxQueue->uxMessagesWaiting > ( unsigned portBASE_TYPE ) 0 )
698     {
699         /* Copy the data from the queue. */
700         pxQueue->pcReadFrom += pxQueue->uxItemSize;
701         if( pxQueue->pcReadFrom >= pxQueue->pcTail )
702         {
703             pxQueue->pcReadFrom = pxQueue->pcHead;
704         }
705         --( pxQueue->uxMessagesWaiting );
706         memcpy( ( void * ) pvBuffer, ( void * ) pxQueue->pcReadFrom,
707                 ( unsigned ) pxQueue->uxItemSize );
708
709         /* If the queue is locked we will not modify the event list. Instead we update the lock count
710         so the task that unlocks the queue will know that an ISR has removed data while the queue was
711         locked. */
712         if( pxQueue->xRxLock == queueUNLOCKED )
713         {
714             /* We only want to wake one task per ISR, so check that a task has not already been woken. */
715             if( !( *pxTaskWoken ) )
716             {
717                 if( !listLIST_IS_EMPTY( &( pxQueue->xTasksWaitingToSend ) ) )
718                 {
719                     if( xTaskRemoveFromEventList( &( pxQueue->xTasksWaitingToSend ) )
720                         != pdFALSE )
721                     {
722                         /* The task waiting has a higher priority than us so
723                         force a context switch. */
724                         *pxTaskWoken = pdTRUE;
725                     }
726                 }
727             }
728         }
729     }
730     {
731         /* Increment the lock count so the task that unlocks the queue
732         knows that data was removed while it was locked. */
733         ++( pxQueue->xRxLock );
734     }
735
736     xReturn = pdPASS;
737 }
738 else
739 {
740     xReturn = pdFAIL;
741 }
742
743 return xReturn;
744 }
745
746 unsigned portBASE_TYPE uxQueueMessagesWaiting( xQueueHandle pxQueue )
747 {
748     unsigned portBASE_TYPE uxReturn;
749
750     taskENTER_CRITICAL();
751     uxReturn = pxQueue->uxMessagesWaiting;
752     taskEXIT_CRITICAL();
753
754     return uxReturn;
755 }
756
757 void vQueueDelete( xQueueHandle pxQueue )
758 {
759     vPortFree( pxQueue->pcHead );
760     vPortFree( pxQueue );
761 }
762

```

```

763 static signed portBASE_TYPE prvUnlockQueue( xQueueHandle pxQueue )
764 {
765     signed portBASE_TYPE xYieldRequired = pdFALSE;
766
767     /* THIS FUNCTION MUST BE CALLED WITH THE SCHEDULER SUSPENDED. */
768
769     /* The lock counts contains the number of extra data items placed or
770      removed from the queue while the queue was locked. When a queue is
771      locked items can be added or removed, but the event lists cannot be
772      updated. */
773     taskENTER_CRITICAL();
774     {
775         --( pxQueue->xTxLock );
776
777         /* See if data was added to the queue while it was locked. */
778         if( pxQueue->xTxLock > queueUNLOCKED )
779         {
780             pxQueue->xTxLock = queueUNLOCKED;
781
782             /* Data was posted while the queue was locked. Are any tasks
783              blocked waiting for data to become available? */
784             if( !listLIST_IS_EMPTY( &( pxQueue->xTasksWaitingToReceive ) ) )
785             {
786                 /* Tasks that are removed from the event list will get added to
787                  the pending ready list as the scheduler is still suspended. */
788                 if( xTaskRemoveFromEventList( &( pxQueue->xTasksWaitingToReceive ) ) != pdFALSE
789             )
790             {
791                 /* The task waiting has a higher priority so record that a
792                  context switch is required. */
793                 xYieldRequired = pdTRUE;
794             }
795         }
796     }
797 }
798 taskEXIT_CRITICAL();
799
800 /* Do the same for the Rx lock. */
801 taskENTER_CRITICAL();
802 {
803     --( pxQueue->xRxLock );
804
805     if( pxQueue->xRxLock > queueUNLOCKED )
806     {
807         pxQueue->xRxLock = queueUNLOCKED;
808
809         if( !listLIST_IS_EMPTY( &( pxQueue->xTasksWaitingToSend ) ) )
810         {
811             if( xTaskRemoveFromEventList( &( pxQueue->xTasksWaitingToSend ) ) != pdFALSE )
812             {
813                 xYieldRequired = pdTRUE;
814             }
815         }
816     }
817 }
818 taskEXIT_CRITICAL();
819
820 return xYieldRequired;
821 }
822
823 static signed portBASE_TYPE prvIsQueueEmpty( const xQueueHandle pxQueue )
824 {
825     signed portBASE_TYPE xReturn;
826
827     taskENTER_CRITICAL();
828     xReturn = ( pxQueue->uxMessagesWaiting == ( unsigned portBASE_TYPE ) 0 );
829     taskEXIT_CRITICAL();
830
831     return xReturn;
832 }
833
834 static signed portBASE_TYPE prvIsQueueFull( const xQueueHandle pxQueue )
835 {
836     signed portBASE_TYPE xReturn;
837
838     taskENTER_CRITICAL();
839     xReturn = ( pxQueue->uxMessagesWaiting == pxQueue->uxLength );
840     taskEXIT_CRITICAL();
841
842     return xReturn;
843 }

```

Answer to Question 1

1.

a)

The job model assumes that each task starts executing at fixed periodic times (known period), executes for fixed CPU time, then suspends until next period.

The RMA utilisation limit applies when:

Tasks are priority scheduled with priorities such that smaller periods have higher priority

Each task has deadline equal to its period

There is no blocking

[4]

b) Priorities: A>B>D>C (2 marks)

Utilisation: $10+5+2+4 = 21$ cycles/us. Limit is $4(2^{1/4}-1) = 0.757$

so $21/f < 0.757$ (f in cycles/us = MHz) $\Rightarrow f > 27.7$ MHz, $T < 35.1$ ns (3 marks)

[5]

c) Blocking only matters from lower priority task, so:

A: R

B: R

D: R

C: 0

New utilisation adds $R/10+R/40+R/100 = 0.135R$ so

$f > (21 + 0.135R)/0.757$, $T < 0.757/(21+0.135R)$

[5]

d) In the new system D,C are equal priority. This will help C, but slow D. Hence the maximum blocking time (for D) is the execution time of C, and no other task is blocked. Extended RMA gives us $f > (21+3)/0.757 = 31.7$ MHz, $T < 31.5$ ns. [-3 marks if assume incorrectly C & D each block the other, $T < 28.4$ ns]

[6]

Answer to Question 2

a) Four pointers per list node

Insertion at head or tail/removal all O(1)

Insertion into sorted list O(n)

No special cases for start/end of list

used: ready list, event lists, delayed task lists

[8]

b) Sentinel node is used so that zero length lists are not a special cases

[4]

c)

Ready list: This is arranged as array of lists indexed by priority so list is unnecessary except when equal priority tasks exist. However in singly-linked list case deletion from list would be difficult O(n). Insertion slightly faster since two less pointers to update. Other operations similar.

Delayed task lists: sorted list insertion is slightly more slower since must either keep two pointers or use pointer with extra level of indirection.
Removal more difficult O(n) not O(1).

Event lists: as for delayed task lists, same operations needed.

[8]

Answer to Question 3

3.

- a) A queue with n messages is equivalent to a semaphore of token count n . QueueSend, QueueReceive take the place of SemaGive, SemaTake. Initialisation by posting a number of initial messages can therefore implement a binary or counting semaphore.

[4]

- b) There are three potential liveness problems: deadlock between tasks A & B which claim q[1] & q[2] in different orders, and priority inversion on q[1] where any of tasks A,B,C could be delayed by task D, and starvation on q1 for A,B,C.

Conditions for deadlock are specific execution order - it can happen regardless of task priorities. Conditions for priority inversion are that some task (e.g. D) is intermediate in priority between a given pair of tasks of A,B,C, and that this task preempts the lower priority of the pair while it has claimed p[1]. Conditions for starvation are utilisation of q1 by highest priority two of A,B,C is $> 100\%$

[16]

Answer to Question 4

4.

a)

Interrupt locking (critical sections) disables interrupts, and also therefore other tasks, from preemption.

Scheduler locking, allows interrupts, but other tasks cannot pre-empt

Queue locking, allows interrupts to occur, and add or remove data from queues, however state of tasks waiting on queues cannot be changed until queue is unlocked.

[5]

b) See exam notes.

line	interrupt	scheduler	queue	notes
623-626		X	X	allow queue event lists to be changed
635	X	X	X	lock interrupts to allow queue to be unlocked
637-638	X			unlock queue, then scheduler. Can't be preempted because interrupts are locked
640	(X)			Yield.. If task changes interrupts will be enabled for duration of change
643-644	X	X	X	lock queues & scheduler first to allow interrupt unlock
646		X	X	unlock interrupts to minimise latency
650	X	X	X	lock interrupts again, to post message
673		X	X	unlock interrupts as soon as possible
676-678				unlock queues & scheduler since all finished

[15]

ANSWERS

Answer to Question 5

5. Table 5.1 Lists the I/O devices serviced by a microcontroller, together with the required response time, frequency at which the device needs service, and length of service routine. Assume that interrupt entry and return time is E us & R us respectively, and that two interrupt levels are available $P1$ and $P2$ with $P1$ higher priority than $P2$ so that each device may raise a $P1$ or $P2$ interrupt, or no interrupt. Each interrupt has a separate vector, however when more than one device can raise an interrupt both use the same vector.

- a) Every device can be serviced in foreground by one of the two interrupts or in a background polling loop.

The shortest service frequency (100us) is longer than the longest service time (90us). thus if deadlines are met at most one service request from each device during response time.

[5]

- b) The time before start of service A, B is $B_S + R + E, R + E + A_S$

[5]

- c) Sum of service times of other devices at own or higher priority must be less than latency. Disk & network must be higher priority than anything else, but can be in either order. Printer and keyboard can be any order if lower than disk & network. So (assuming we use all interrupts & background):

	Background Loop	P2	P1
1	K	P	D & N
2	K&P	D	N
3	K&P	N	D
4	P	K	D&N

[5]

d)

Strategy 1: $D_L > E + R + N_S$,

(do not matter) $N_L > E + R + D_S, P_L > 3E + 2R + D_S + N_S, K_L > 4E + 3R + P_S + D_S + N_S$

Strategy 2: $N_L > E, D_L > 2E + R + N_S$,

(do not matter) $P_L > 3E + 2R + D_S + N_S + K_S$, (also swap P,K)

Strategy 3: as 2 but swap N,D

Strategy 4: as 1 but swap K,P

[5]

Table 5.2

Answer to Question 6

6.

a)

For each channel: sample_input, FFT, average (total 12 tasks). Average & FFT tasks could be combined, in which case there are 8 tasks.

Keyboard processing, LCD control, serial_out (total 3 tasks).

[6]

b) Let time to process 4096 samples on one channel be T.

$$T > 260 + 80\text{us} \Rightarrow T > 340$$

CPU utilisation for each task is $80/T$. Assuming that FFT processing is the only significant CPU task, and that each channel is independent therefore RMA must be used with 11 tasks as per part a, and that tasks do not block, we have:

$$(80/T)*4 < 0.715 \text{ (RMA for 11 tasks)} \Rightarrow T > 448\text{us}$$

The FFT unit itself has utilisation $4*340/T < 1$. This is clearly the limiting factor. Hence max sample rate is for $T > 1360\text{us} \Rightarrow 3.01 \text{ samples/us}$.

Hence max sample rate is 9.14 samples/us. (Answers which estimate blocking are also allowed).

[8]

c)

The FFT is a single resource with mutually exclusive hardware access. It should be protected by a binary semaphore. In this system FFT unit utilisation is much more significant than CPU utilisation, so although priority inversion is possible it will not affect performance.

[6]