

IMPERIAL COLLEGE LONDON

DEPARTMENT OF ELECTRICAL AND ELECTRONIC ENGINEERING
EXAMINATIONS 2014

MSc and EEE PART IV and EIE PART III: MEng and ACGI

Corrected Copy

EMBEDDED SYSTEMS

Friday, 16 May 2:30 pm

Time allowed: 1:30 hours

There are TWO questions on this paper.

Answer BOTH questions.

The questions carry equal marks.

Any special instructions for invigilators and information for candidates are on page 1.

Examiners responsible	First Marker(s) :	T.J.W. Clarke
	Second Marker(s) :	Y.K. Demiris

The Questions

1. An ARM Cortex processor clocked at 50MHz interfaces with a custom real time clock (RTC) module. Time is provided as the values of 4 32 bit registers memory mapped at offsets (0,4,8,12) from address EXT_RTC_BASE. Each register holds 16 bits of the double-word time value in its least significant half-word and zero in its most significant half-word. Registers with lower offsets are less significant. The clock module hardware increments its internal registers synchronously every ARM clock cycle, or 20ns.
- a) Write a C function GetRTC32() without critical sections that in absence of interrupts is guaranteed to return the correct least significant 32 bits of the time value. Explain what is the data consistency problem in this operation, and how your code overcomes it. [5]
- b) The error in the function GetRTC32() is defined to be the difference between the returned value of time, and the time at the end of the return from the function. If there is an interrupt of period 1ms with ISR of length 100 clock cycles including context switch overhead, estimate, with reasons, the largest and smallest value of this error. [4]
- c) Now suppose there are two similar asynchronous interrupts both of period approximately 10 μ s and length, not counting context switch overhead, of 80 cycles each. What is the maximum error as defined in the last part, assuming that FPU state is not saved on exceptions, and ignoring any effects from interrupt late arrival and tail chaining optimisations in the NVIC? [3]
- d) Detail, giving reasons for your answers, for the case described in part c, how are the maximum and the average error affected by:
- i) Late arrival optimisation in the NVIC. [2]
- ii) Tail chaining optimisation in the NVIC. [2]
- e) In multi-processor system a task on one CPU reads the 64 bit RTC module time as in part a and writes this value at intervals into a double-word time[0], time[1] in shared memory. This value is read by tasks running on other CPUs. Both reading and writing is least significant word first. Provide one or more execution traces to illustrate all the distinct ways in which this memory read operation could result in inconsistent data read, and state what is the possible error and the conditions under which such a trace might happen. [4]

2. Figure 2.1 details the CPU time and period of a set of 4 tasks in a real-time system. The tasks are specified using a job model in which new work of the specified CPU time is available periodically with the given fixed period. Each task has the hard deadline that it must complete its work before the end of its period. CPU work is specified in terms of CPU clock cycles, and all context switch and semaphore API overheads should be ignored. The minimum CPU clock frequency at which the system is guaranteed to be schedulable, f_0 , is calculated below under various different conditions.

a) Suppose that there is no blocking. Using Rate Monotonic Analysis (RMA) determine f_0 . Specify the relative priority of tasks required for this guarantee.

[5]

b) Suppose all tasks share a resource used once per task period. Task priorities are as determined by RMA above. The resource access is protected by a binary semaphore that introduces blocking. The CPU time using the resource is 10 cycles, and included in the time shown in Figure 2.1. Calculate the maximum extra time (in clock cycles) for each task due to semaphore blocking assuming that each task blocks any other task at most once. State which tasks suffer priority inversion.

[4]

c) Repeat the analysis in part b assuming that a mutex implementing Priority Inheritance Protocol protects the resource.

[4]

d) Determine f_0 in this system in each of the binary semaphore and mutex cases.

[2]

e) State the potential advantages and disadvantages of using Priority Threshold Scheduling (PTS) in the system of part a. Suppose all tasks are implemented with no pre-emption using co-routines and priorities as determined by RMA. Determine the maximum delay in start time of each task due to lack of preemption. Therefore calculate f_0 in this case. How can this no pre-emption system be modelled using PTS?

[5]

Task	CPU time clocks	Period μs
A	1000	1000
B	200	500
C	100	2000
D	50	100

Figure 2.1: Task Specification.